

Project- CCPS 844 Data Mining

Salman AlMaskati

This project aims to tackle a challenging regression problem focused on predicting life expectancy. Life expectancy is a critical indicator of the overall health and well-being of a population, and accurately predicting it can have far-reaching implications for public health, social policies, and resource allocation. \ \ The key steps of the project involve data cleaing, visualizations, Clustering, feature selection and PCA, model training and evaluation. Several regression learing algorithms will be explored and compared to find the best-performing model.

Table of Contents

Data Summary

Web Scrapping

Data Preprocessing

Visualizations

EDA, Clustering

Feature Elimination (RFECV)

PCA

Spliting Data

Regression

Regression algorithms Summary

Report

Data summary

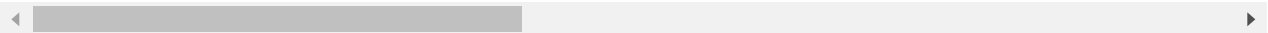
```
In [ ]: import pandas as pd

df=pd.read_csv("world-data-2023.csv")
df
```

Out[]:

	Country	Density\n(P/Km2)	Abbreviation	Agricultural Land(%)	Land Area(Km2)	Armed Forces size	Birth Rate	Calling Code	Capi
0	Afghanistan	60	AF	58.10%	652,230	323,000	32.49	93.0	
1	Albania	105	AL	43.10%	28,748	9,000	11.78	355.0	
2	Algeria	18	DZ	17.40%	2,381,741	317,000	24.28	213.0	
3	Andorra	164	AD	40.00%	468	NaN	7.20	376.0	
4	Angola	26	AO	47.50%	1,246,700	117,000	40.73	244.0	
...	
190	Venezuela	32	VE	24.50%	912,050	343,000	17.88	58.0	
191	Vietnam	314	VN	39.30%	331,210	522,000	16.75	84.0	
192	Yemen	56	YE	44.60%	527,968	40,000	30.45	967.0	
193	Zambia	25	ZM	32.10%	752,618	16,000	36.19	260.0	
194	Zimbabwe	38	ZW	41.90%	390,757	51,000	30.68	263.0	

195 rows × 35 columns



```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 195 entries, 0 to 194
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Country                               195 non-null    object
1   Density                               195 non-null    object
   (P/Km2)
2   Abbreviation                         188 non-null    object
3   Agricultural Land( %)                188 non-null    object
4   Land Area(Km2)                       194 non-null    object
5   Armed Forces size                    171 non-null    object
6   Birth Rate                           189 non-null    float64
7   Calling Code                         194 non-null    float64
8   Capital/Major City                   192 non-null    object
9   Co2-Emissions                        188 non-null    object
10  CPI                                  178 non-null    object
11  CPI Change (%)                       179 non-null    object
12  Currency-Code                        180 non-null    object
13  Fertility Rate                       188 non-null    float64
```

```

14 Forested Area (%) 188 non-null object
15 Gasoline Price 175 non-null object
16 GDP 193 non-null object
17 Gross primary education enrollment (%) 188 non-null object
18 Gross tertiary education enrollment (%) 183 non-null object
19 Infant mortality 189 non-null float64
20 Largest city 189 non-null object
21 Life expectancy 187 non-null float64
22 Maternal mortality ratio 181 non-null float64
23 Minimum wage 150 non-null object
24 Official language 194 non-null object
25 Out of pocket health expenditure 188 non-null object
26 Physicians per thousand 188 non-null float64
27 Population 194 non-null object
28 Population: Labor force participation (%) 176 non-null object
29 Tax revenue (%) 169 non-null object
30 Total tax rate 183 non-null object
31 Unemployment rate 176 non-null object
32 Urban_population 190 non-null object
33 Latitude 194 non-null float64
34 Longitude 194 non-null float64

```

dtypes: float64(9), object(26)

memory usage: 53.4+ KB

In []:

```
df.describe()
```

Out[]:

	Birth Rate	Calling Code	Fertility Rate	Infant mortality	Life expectancy	Maternal mortality ratio	Physicians per thousand	Latitude
count	189.000000	194.000000	188.000000	189.000000	187.000000	181.000000	188.000000	194.000000
mean	20.214974	360.546392	2.698138	21.332804	72.279679	160.392265	1.839840	19.092351
std	9.945774	323.236419	1.282267	19.548058	7.483661	233.502024	1.684261	23.961779
min	5.900000	1.000000	0.980000	1.400000	52.800000	2.000000	0.010000	-40.900557
25%	11.300000	82.500000	1.705000	6.000000	67.000000	13.000000	0.332500	4.544175
50%	17.950000	255.500000	2.245000	14.000000	73.200000	53.000000	1.460000	17.273849
75%	28.750000	506.750000	3.597500	32.700000	77.500000	186.000000	2.935000	40.124603
max	46.080000	1876.000000	6.910000	84.500000	85.400000	1150.000000	8.420000	64.963051

Dropping columns that do not add value to the df

In []:

```

cols_to_drop= ['Abbreviation', 'Calling Code', 'Capital/Major City', 'Currency-Code', 'Large
df = df.drop(columns=cols_to_drop)
df

```

Out[]:

	Country	Density\n(P/Km2)	Agricultural Land(%)	Land Area(Km2)	Armed Forces size	Birth Rate	Co2-Emissions	CPI	CF Chang (%)
0	Afghanistan	60	58.10%	652,230	323,000	32.49	8,672	149.9	2.30%

	Country	Density\n(P/Km2)	Agricultural Land(%)	Land Area(Km2)	Armed Forces size	Birth Rate	Co2-Emissions	CPI	CF Chang (%)
1	Albania	105	43.10%	28,748	9,000	11.78	4,536	119.05	1.40%
2	Algeria	18	17.40%	2,381,741	317,000	24.28	150,006	151.36	2.00%
3	Andorra	164	40.00%	468	NaN	7.20	469	NaN	NaN
4	Angola	26	47.50%	1,246,700	117,000	40.73	34,693	261.73	17.10%
...
190	Venezuela	32	24.50%	912,050	343,000	17.88	164,175	2,740.27	254.90%
191	Vietnam	314	39.30%	331,210	522,000	16.75	192,668	163.52	2.80%
192	Yemen	56	44.60%	527,968	40,000	30.45	10,609	157.58	8.10%
193	Zambia	25	32.10%	752,618	16,000	36.19	5,141	212.31	9.20%
194	Zimbabwe	38	41.90%	390,757	51,000	30.68	10,983	105.51	0.90%

195 rows × 24 columns

In []: `df.shape`

Out[]: (195, 24)

Visualizing NA's

```
In [ ]: import seaborn as sns
import matplotlib.pyplot as plt

na_counts = df.isna().sum()
print("The Number of NA's in",na_counts)

sns.heatmap(df.isnull(),yticklabels=False,cbar=False)
plt.title("Visualizing NA's")
```

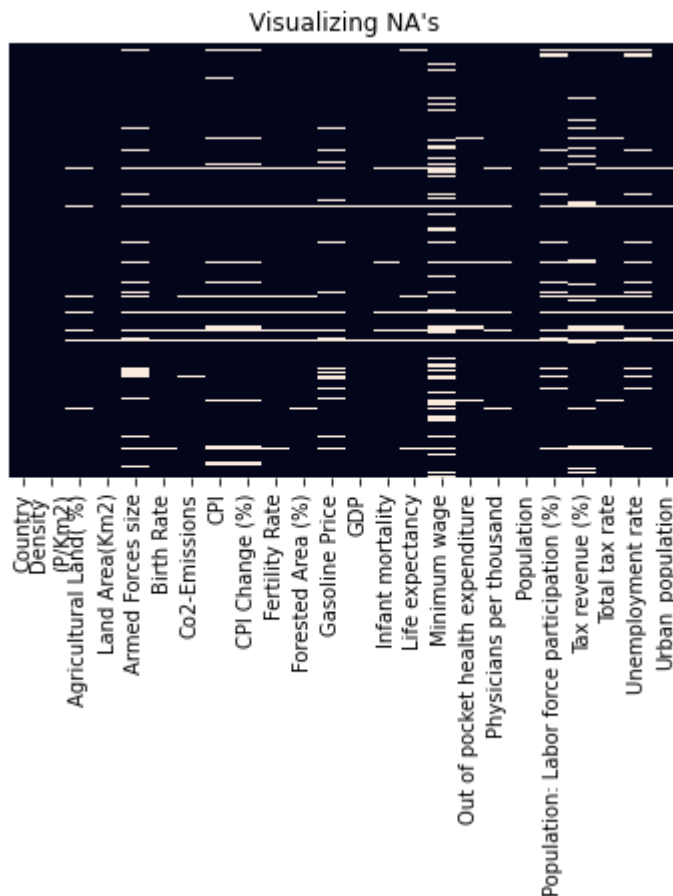
```
The Number of NA's in Country          0
Density\n(P/Km2)                      0
Agricultural Land( %)                  7
Land Area(Km2)                        1
Armed Forces size                     24
Birth Rate                            6
Co2-Emissions                         7
CPI                                   17
CPI Change (%)                       16
Fertility Rate                        7
Forested Area (%)                    7
Gasoline Price                       20
GDP                                   2
Infant mortality                      6
Life expectancy                      8
Minimum wage                         45
```

```

Out of pocket health expenditure    7
Physicians per thousand             7
Population                          1
Population: Labor force participation (%) 19
Tax revenue (%)                     26
Total tax rate                      12
Unemployment rate                   19
Urban_population                    5
dtype: int64

```

```
Out[ ]: Text(0.5, 1.0, "Visualizing NA's")
```



Web Scraping

Fill in Agricultural Land(%) values

```

In [ ]: from lxml import html
        from bs4 import BeautifulSoup
        import requests

```

```

In [ ]: #Find Out which countries have missing Agricultural Land( %) values
        rows_with_na = df[df["Agricultural Land( %)"].isna()]

        # Display the result
        print(rows_with_na)

        print("Eswatini, Monaco, Nauru, North Macedonia, Palestinian National Authority, South

```

	Country	Density\n(P/Km2)	Agricultural Land(%)	\
56	Eswatini	67	NaN	
73	Vatican City	2,003	NaN	
113	Monaco	26,337	NaN	
120	Nauru	541	NaN	
128	North Macedonia	83	NaN	
133	Palestinian National Authority	847	NaN	
163	South Sudan	18	NaN	

	Land Area(Km2)	Armed Forces size	Birth Rate	Co2-Emissions	CPI	\
56	17,364	NaN	NaN	NaN	NaN	
73	0	NaN	NaN	NaN	NaN	
113	2	NaN	5.90	NaN	NaN	
120	21	NaN	NaN	NaN	NaN	
128	25,713	NaN	NaN	NaN	NaN	
133	NaN	NaN	NaN	NaN	NaN	
163	644,329	185,000	35.01	1,727	4,583.71	

	CPI Change (%)	Fertility Rate	... Life expectancy	Minimum wage	\
56	NaN	NaN	...	NaN	NaN
73	NaN	NaN	...	NaN	NaN
113	NaN	NaN	...	NaN	\$11.72
120	NaN	NaN	...	NaN	NaN
128	NaN	NaN	...	NaN	NaN
133	NaN	NaN	...	NaN	NaN
163	187.90%	4.7	...	57.6	NaN

	Out of pocket health expenditure	Physicians per thousand	Population	\
56	11.30%	NaN	1,093,238	
73	NaN	NaN	836	
113	6.10%	6.56	38,964	
120	NaN	NaN	10,084	
128	35.60%	NaN	1,836,713	
133	NaN	NaN	NaN	
163	61.30%	NaN	11,062,113	

	Population: Labor force participation (%)	Tax revenue (%)	Total tax rate	\
56	NaN	28.60%	NaN	
73	NaN	NaN	NaN	
113	NaN	NaN	NaN	
120	NaN	NaN	NaN	
128	NaN	NaN	NaN	
133	NaN	NaN	NaN	
163	72.40%	NaN	31.40%	

	Unemployment rate	Urban_population
56	NaN	NaN
73	NaN	NaN
113	NaN	38,964
120	NaN	NaN
128	NaN	NaN
133	NaN	NaN
163	12.24%	2,201,250

[7 rows x 24 columns]

Eswatini, Monaco, Nauru, North Macedonia, Palestinian National Authority, South Sudan

```
In [ ]: #get agricultral land data
import requests
```

```
from bs4 import BeautifulSoup
```

In []:

```
url = 'https://wdi.worldbank.org/table/3.2'

# end a GET request to the URL
response = requests.get(url)
html_content = response.text

# create a BeautifulSoup object
soup = BeautifulSoup(html_content, 'html.parser')

country_elements = soup.find_all(class_="country")

#Eswatini %
for country_element in country_elements:
    country_content = country_element.get_text()

    if "Eswatini" in country_content:
        row_element = country_element.find_parent("tr")
        if row_element:
            eswatini_row = row_element.get_text()
            #print("Row:", eswatini_row)

#Monaco %
for country_element in country_elements:
    country_content = country_element.get_text()

    if "Monaco" in country_content:
        row_element = country_element.find_parent("tr")
        if row_element:
            monaco_row = row_element.get_text()
            #print("Row:", monaco_row)

#Nauru %
for country_element in country_elements:
    country_content = country_element.get_text()

    if "North Macedonia" in country_content:
        row_element = country_element.find_parent("tr")
        if row_element:
            nm_row = row_element.get_text()
            #print("Row:", nm_row)

#South Sudan %
for country_element in country_elements:
    country_content = country_element.get_text()

    if "South Sudan" in country_content:
        row_element = country_element.find_parent("tr")
        if row_element:
            ss_row = row_element.get_text()
            #print("Row:", ss_row)

print("Nauru, Vatican City, and Moanco, could not be found\nPalestinian National Author

# helper function to format extaracted data
def get_first_number(words):
    for word in words:
```

```

        if word.replace('.', '', 1).isdigit(): # forming
            return float(word)

# split the rows by whitespaces
nm_words = nm_row.split()
ss_words = ss_row.split()
eswatini_words = eswatini_row.split()

# create dictionary
data_dict = {
    'North Macedonia': get_first_number(nm_words),
    'South Sudan': get_first_number(ss_words),
    'Eswatini': get_first_number(eswatini_words)
}

print(data_dict)

```

Nauru, Vatican City, and Moanco, could not be found
 Palestinian National Authority(will be dropped due to lack of data)
 {'North Macedonia': 50.0, 'South Sudan': 45.0, 'Eswatini': 71.0}

```

In [ ]: #input values to df
df.loc[df['Country'] == 'North Macedonia', 'Agricultural Land( %)'] = get_first_number(
df.loc[df['Country'] == 'South Sudan', 'Agricultural Land( %)'] = get_first_number(ss_w
df.loc[df['Country'] == 'Eswatini', 'Agricultural Land( %)'] = get_first_number(eswatin

```

Data Preprocessing

Replace '0' in all rows with NA

```

In [ ]: import numpy as np
df = df.replace(0, np.nan)

```

```

In [ ]: sns.heatmap(df.isnull(),yticklabels=False,cbar=False)
plt.title("Visualizing NA's")
na_counts = df.isna().sum()
print("The Number of NA's in",na_counts)

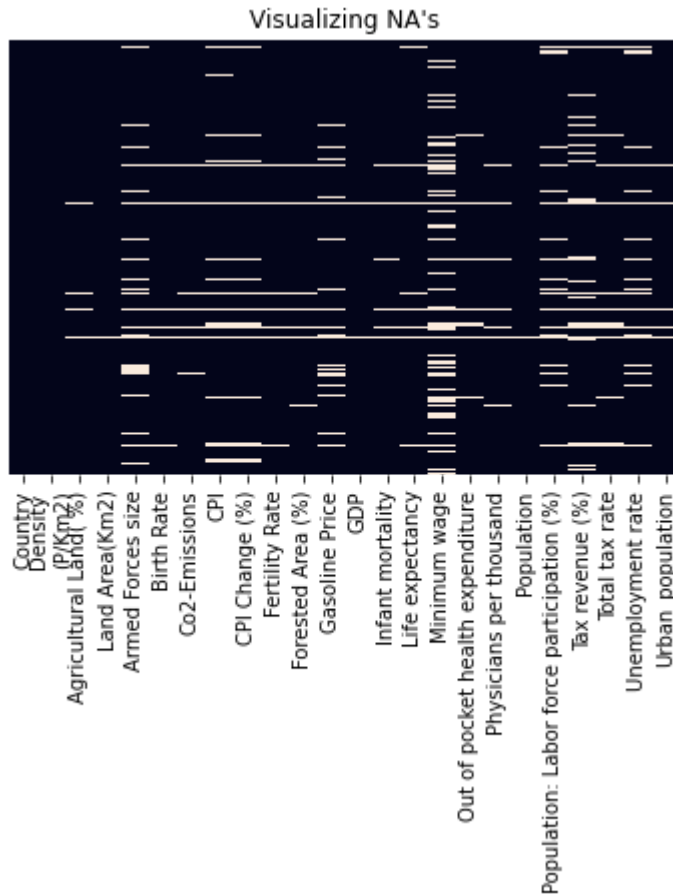
```

The Number of NA's in Country	0
Density\n(P/Km2)	0
Agricultural Land(%)	4
Land Area(Km2)	1
Armed Forces size	24
Birth Rate	6
Co2-Emissions	7
CPI	17
CPI Change (%)	16
Fertility Rate	7
Forested Area (%)	7
Gasoline Price	20
GDP	2
Infant mortality	6
Life expectancy	8
Minimum wage	45


```

Out of pocket health expenditure    7
Physicians per thousand            7
Population                        1
Population: Labor force participation (%) 19
Tax revenue (%)                   26
Total tax rate                    12
Unemployment rate                 19
Urban_population                   5
dtype: int64

```



Drop Minumum Wage column due to lack of data

```
In [ ]: df.drop(columns='Minimum wage', inplace=True)
```

Drop ROWS that have NA values

```
In [ ]: df = df.dropna(axis=0)
```

```
In [ ]: print(df.isnull().sum())
```

```

Country                0
Density\n(P/Km2)       0
Agricultural Land( %)  0
Land Area(Km2)         0
Armed Forces size      0
Birth Rate             0
Co2-Emissions          0
CPI                   0
CPI Change ( %)       0

```

```

Fertility Rate      0
Forested Area (%)  0
Gasoline Price      0
GDP                 0
Infant mortality    0
Life expectancy     0
Out of pocket health expenditure  0
Physicians per thousand  0
Population          0
Population: Labor force participation (%)  0
Tax revenue (%)     0
Total tax rate      0
Unemployment rate   0
Urban_population    0
dtype: int64

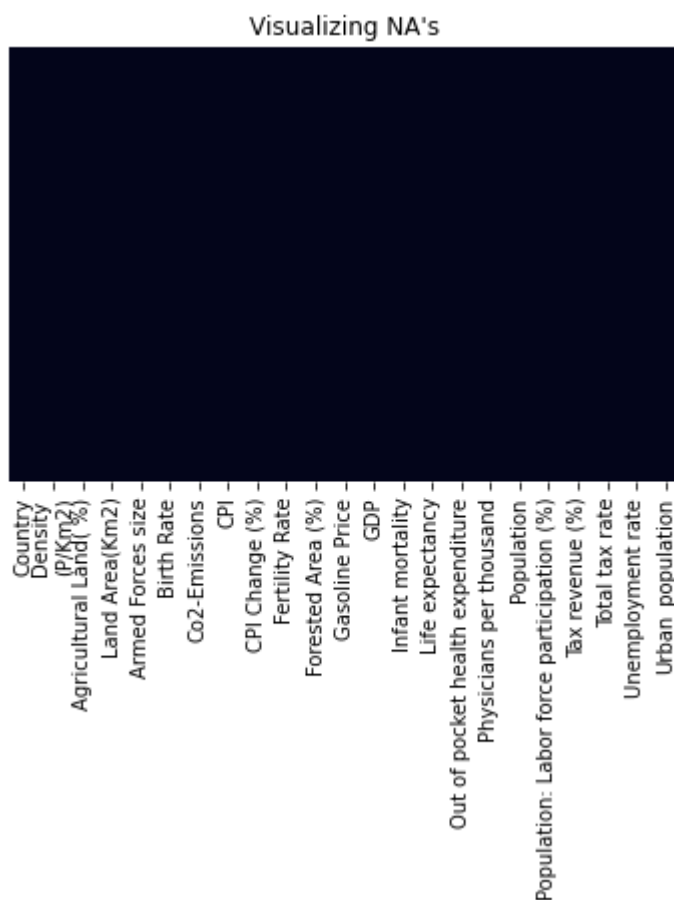
```

```
In [ ]: df.shape
```

```
Out[ ]: (146, 23)
```

```
In [ ]: sns.heatmap(df.isnull(),yticklabels=False,cbar=False)
plt.title("Visualizing NA's")
```

```
Out[ ]: Text(0.5, 1.0, "Visualizing NA's")
```



change data types 'object' to 'float64' so we can fit the dataframe into machine learning algorithms

In []: `df.dtypes`

```
Out[ ]: Country                object
Density\n(P/Km2)            object
Agricultural Land( %)      object
Land Area(Km2)              object
Armed Forces size           object
Birth Rate                  float64
Co2-Emissions               object
CPI                         object
CPI Change (%)              object
Fertility Rate              float64
Forested Area (%)           object
Gasoline Price              object
GDP                         object
Infant mortality            float64
Life expectancy             float64
Out of pocket health expenditure object
Physicians per thousand     float64
Population                  object
Population: Labor force participation (%) object
Tax revenue (%)             object
Total tax rate              object
Unemployment rate           object
Urban_population            object
dtype: object
```

```
In [ ]: #Changing columns data types

#remove commas from Density column
df['Density\n(P/Km2)'] = df['Density\n(P/Km2)'].str.replace(',', '')
df['Density\n(P/Km2)'] = df['Density\n(P/Km2)'].astype(float)

#remove (%) from Agricultural Land(%) column
df['Agricultural Land( %)'] = df['Agricultural Land( %)'].str.replace('%', '')
df['Agricultural Land( %)'] = df['Agricultural Land( %)'].astype(float)

#remove commas from Land Area(Km2) column
df['Land Area(Km2)'] = df['Land Area(Km2)'].str.replace(',', '')
df['Land Area(Km2)'] = df['Land Area(Km2)'].astype(float)

#remove commas from Armed Forces size column
df['Armed Forces size'] = df['Armed Forces size'].str.replace(',', '')
df['Armed Forces size'] = df['Armed Forces size'].astype(float)

#remove commas from Co2-Emissions column
df['Co2-Emissions'] = df['Co2-Emissions'].str.replace(',', '')
df['Co2-Emissions'] = df['Co2-Emissions'].astype(float)

#remove commas from CPI column
df['CPI'] = df['CPI'].str.replace(',', '')
df['CPI'] = df['CPI'].astype(float)

#remove "%" from CPI Change column
df['CPI Change (%)'] = df['CPI Change (%)'].str.replace('%', '')
df['CPI Change (%)'] = df['CPI Change (%)'].astype(float)
```

```

#remove "%" from Forested Area (%) column
df['Forested Area (%)'] = df['Forested Area (%)'].str.replace('%', '')
df['Forested Area (%)'] = df['Forested Area (%)'].astype(float)

#remove "$" from Gasoline Price column
df['Gasoline Price'] = df['Gasoline Price'].str.replace('$', '')
df['Gasoline Price'] = df['Gasoline Price'].astype(float)

#remove "$" from GDP Change column
df['GDP'] = df['GDP'].str.replace('$', '')
#remove commas from GDP Change column
df['GDP'] = df['GDP'].str.replace(',', '')
df['GDP'] = df['GDP'].astype(float)
#Pandas formats large number to scientific notation, therefore this converts back
pd.set_option('display.float_format', '{:.2f}'.format)

#remove "%" from Out of pocket health expenditure column
df['Out of pocket health expenditure'] = df['Out of pocket health expenditure'].str.replace('%', '')
df['Out of pocket health expenditure'] = df['Out of pocket health expenditure'].astype(float)

#remove commas from Population column
df['Population'] = df['Population'].str.replace(',', '')
df['Population'] = df['Population'].astype(float)

#remove "%" from Population: Labor force participation (%) column
df['Population: Labor force participation (%)'] = df['Population: Labor force participation (%)'].str.replace('%', '')
df['Population: Labor force participation (%)'] = df['Population: Labor force participation (%)'].astype(float)

#remove "%" from Tax revenue (%) column
df['Tax revenue (%)'] = df['Tax revenue (%)'].str.replace('%', '')
df['Tax revenue (%)'] = df['Tax revenue (%)'].astype(float)

#remove "%" from Total tax rate column
df['Total tax rate'] = df['Total tax rate'].str.replace('%', '')
df['Total tax rate'] = df['Total tax rate'].astype(float)

#remove "%" from Unemployment rate column
df['Unemployment rate'] = df['Unemployment rate'].str.replace('%', '')
df['Unemployment rate'] = df['Unemployment rate'].astype(float)

#remove commas from Urban population column
df['Urban population'] = df['Urban population'].str.replace(',', '')
df['Urban population'] = df['Urban population'].astype(float)

#turn life expectancy into int so it can be used as a TARGET variable
df['Life expectancy'] = df['Life expectancy'].astype(int)

```

C:\Users\almas\AppData\Local\Temp\ipykernel_8696\1912219860.py:37: FutureWarning: The default value of regex will change from True to False in a future version. In addition, single character regular expressions will *not* be treated as literal strings when regex=True.

```
df['Gasoline Price'] = df['Gasoline Price'].str.replace('$', '')
```

C:\Users\almas\AppData\Local\Temp\ipykernel_8696\1912219860.py:41: FutureWarning: The default value of regex will change from True to False in a future version. In addition, single character regular expressions will *not* be treated as literal strings when regex=True.

```
df['GDP'] = df['GDP'].str.replace('$', '')
```

All NUMERIC columns that were data type 'object' are now changed to 'float64'. Target Variable 'Life Expectancy' is changed into 'int32'

In []:

```
df.dtypes
```

Out[]:

Country	object
Density\n(P/Km2)	float64
Agricultural Land(%)	float64
Land Area(Km2)	float64
Armed Forces size	float64
Birth Rate	float64
Co2-Emissions	float64
CPI	float64
CPI Change (%)	float64
Fertility Rate	float64
Forested Area (%)	float64
Gasoline Price	float64
GDP	float64
Infant mortality	float64
Life expectancy	int32
Out of pocket health expenditure	float64
Physicians per thousand	float64
Population	float64
Population: Labor force participation (%)	float64
Tax revenue (%)	float64
Total tax rate	float64
Unemployment rate	float64
Urban_population	float64
dtype:	object

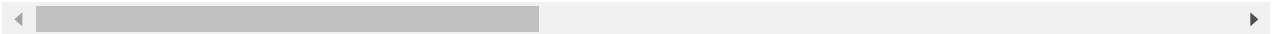
In []:

```
df
```

Out[]:

	Country	Density\n(P/Km2)	Agricultural Land(%)	Land Area(Km2)	Armed Forces size	Birth Rate	Co2-Emissions	CPI	Cha
0	Afghanistan	60.00	58.10	652230.00	323000.00	32.49	8672.00	149.90	
1	Albania	105.00	43.10	28748.00	9000.00	11.78	4536.00	119.05	
2	Algeria	18.00	17.40	2381741.00	317000.00	24.28	150006.00	151.36	
4	Angola	26.00	47.50	1246700.00	117000.00	40.73	34693.00	261.73	1
6	Argentina	17.00	54.30	2780400.00	105000.00	17.02	201348.00	232.75	5
...	
186	United States	36.00	44.40	9833517.00	1359000.00	11.60	5006302.00	117.24	
187	Uruguay	20.00	82.60	176215.00	22000.00	13.86	6766.00	202.92	
191	Vietnam	314.00	39.30	331210.00	522000.00	16.75	192668.00	163.52	
193	Zambia	25.00	32.10	752618.00	16000.00	36.19	5141.00	212.31	
194	Zimbabwe	38.00	41.90	390757.00	51000.00	30.68	10983.00	105.51	

146 rows × 23 columns



Visualizations

Correlation matrix in text and heatmap (to idenitfy relationships and to help with feature selection)

In []:

```
#correlation matrix
correlation_matrix = df.corr()

plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='RdBu_r', center=0, fmt='.2f',
            xticklabels=correlation_matrix.columns, yticklabels=correlation_matrix.colu
plt.title('Correlation Heatmap')
plt.xticks(rotation=90, ha='right')
plt.yticks(rotation=0)

# Convert heatmap to text representaion
correlation_text = correlation_matrix.to_string(float_format="{:.2f}".format)

print("Correlation Heatmap (Text Representation):")
print(correlation_text)

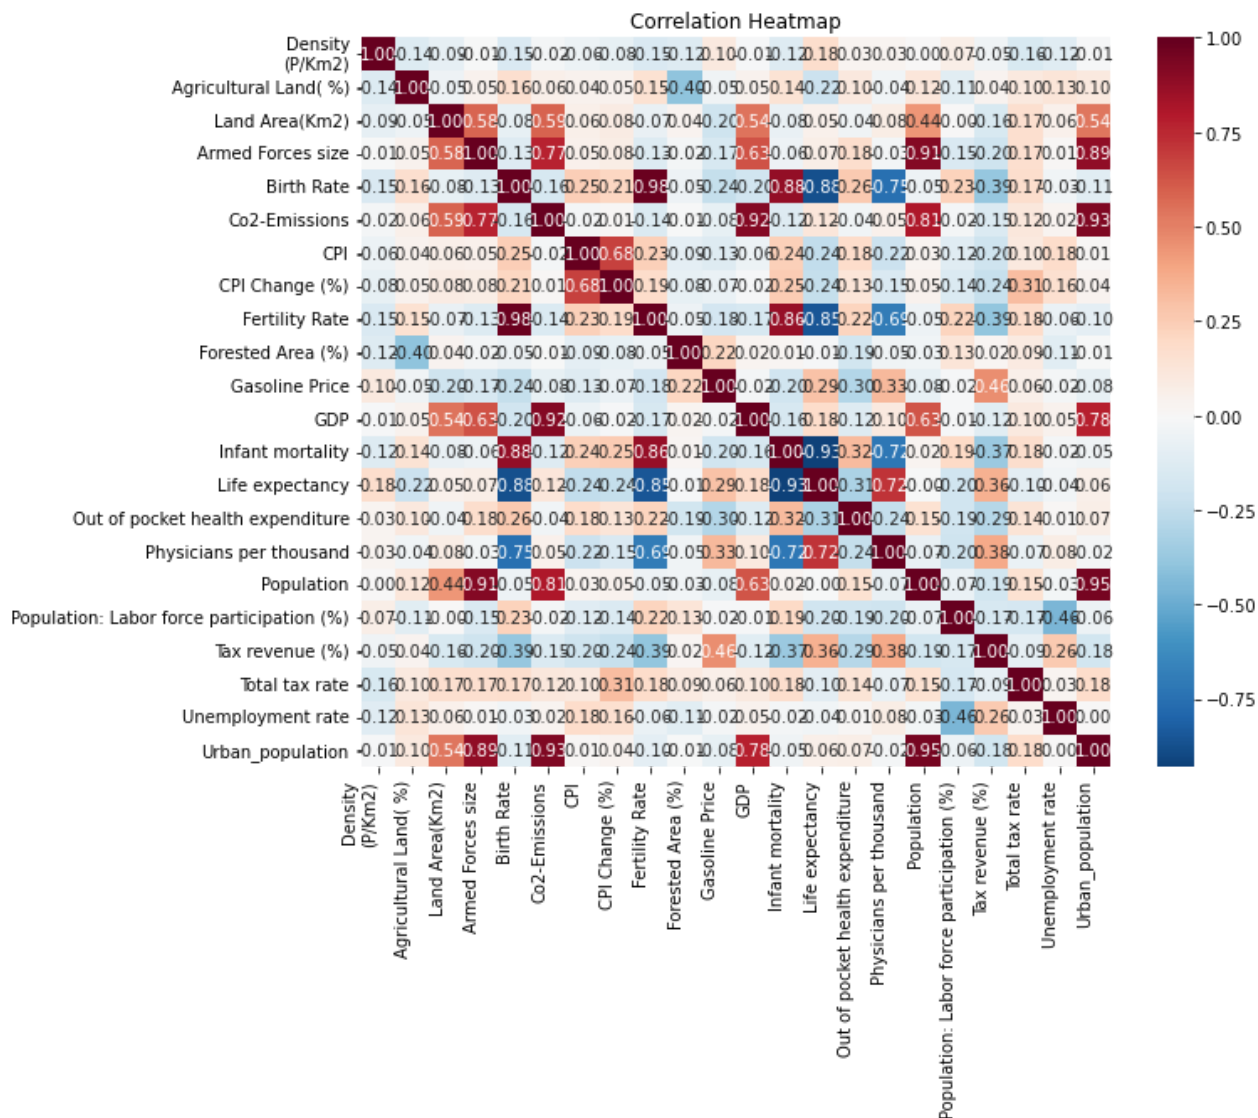
plt.show()
```

Correlation Heatmap (Text Representation):

	Density\n(P/Km2)	Agricultural Land(%)	Land Area(Km2)	Armed Forces size	Birth Rate	Co2-Emissions	CPI	CPI Change (%)	Fertility Rate	Forested Area (%)	Gasoline Price	GDP	Infant mortality	Life expectancy	Out of
--	------------------	-----------------------	----------------	-------------------	------------	---------------	-----	----------------	----------------	-------------------	----------------	-----	------------------	-----------------	--------

pocket health expenditure participation (%)	Physicians per thousand	Population	Population: Labor force
Tax revenue (%)	Total tax rate	Unemployment rate	Urban_population
Density\n(P/Km2)		1.00	-0.14
-0.09	-0.01	-0.15	-0.02
5	-0.12	-0.06	-0.08
0.03	0.03	-0.12	0.18
-0.05	-0.16	-0.12	0.07
Agricultural Land(%)		-0.01	1.00
-0.05	0.05	-0.14	0.05
5	-0.40	0.06	-0.22
0.10	-0.04	0.04	-0.11
0.04	0.10	0.12	0.10
Land Area(Km2)		0.13	-0.09
1.00	0.58	0.59	-0.05
0.04	-0.20	0.06	-0.07
-0.04	0.54	0.05	-0.00
-0.16	0.08	0.44	0.05
Armed Forces size		0.06	-0.01
0.58	1.00	0.54	0.08
-0.02	-0.17	-0.01	-0.13
0.18	0.63	0.07	-0.15
-0.20	-0.03	0.91	0.16
Birth Rate		0.01	0.25
-0.08	0.17	0.89	0.21
8	-0.05	-0.15	-0.88
0.26	-0.75	0.88	0.23
-0.39	0.17	-0.03	-0.11
Co2-Emissions		-0.11	-0.02
0.59	0.77	-0.02	0.06
-0.01	-0.08	0.92	-0.14
-0.04	0.05	-0.12	-0.02
-0.15	0.12	0.81	0.93
CPI		0.02	-0.06
0.06	0.05	0.25	0.04
-0.09	-0.13	-0.02	0.68
0.18	-0.06	0.24	0.23
-0.20	-0.22	0.03	-0.12
CPI Change (%)		0.18	0.01
0.08	0.08	-0.08	0.05
-0.08	-0.07	0.21	1.00
0.13	-0.02	0.25	0.19
-0.24	-0.15	0.05	-0.14
Fertility Rate		0.16	0.04
-0.07	-0.13	-0.15	0.15
0	-0.05	0.23	0.19
0.22	-0.69	0.86	-0.85
-0.39	0.18	-0.05	0.22
Forested Area (%)		-0.06	-0.10
0.04	-0.02	-0.12	-0.40
1.00	0.22	-0.01	-0.09
-0.19	0.02	-0.01	-0.01
0.02	0.09	-0.11	-0.01
Gasoline Price		0.10	-0.05
-0.20	-0.17	-0.24	-0.08
8	0.22	1.00	-0.13
-0.30	0.33	-0.02	-0.20
0.46	0.06	-0.08	-0.08
GDP		-0.02	-0.01
0.54	0.63	-0.20	0.92
0.02	-0.02	-0.16	-0.06

-0.12		0.10	0.63			-0.01
-0.12	0.10		0.05	0.78		
Infant mortality				-0.12	0.14	
-0.08	-0.06	0.88	-0.12	0.24	0.25	0.8
6	0.01	-0.20	-0.16	1.00	-0.93	
0.32	-0.72	0.02			0.19	
-0.37	0.18	-0.02		-0.05		
Life expectancy				0.18	-0.22	
0.05	0.07	-0.88	0.12	-0.24	-0.24	-0.85
-0.01	0.29	0.18	-0.93	1.00		
-0.31		0.72	-0.00		-0.20	
0.36	-0.10	-0.04		0.06		
Out of pocket health expenditure				0.03	0.10	
-0.04	0.18	0.26	-0.04	0.18	0.13	0.2
2	-0.19	-0.30	-0.12	0.32	-0.31	
1.00	-0.24	0.15			-0.19	
-0.29	0.14	0.01		0.07		
Physicians per thousand				0.03	-0.04	
0.08	-0.03	-0.75	0.05	-0.22	-0.15	-0.69
-0.05	0.33	0.10	-0.72	0.72		
-0.24		1.00	-0.07		-0.20	
0.38	-0.07	0.08		-0.02		
Population				0.00	0.12	
0.44	0.91	-0.05	0.81	0.03	0.05	-0.05
-0.03	-0.08	0.63	0.02	-0.00		
0.15	-0.07	1.00			-0.07	
-0.19	0.15	-0.03		0.95		
Population: Labor force participation (%)				0.07	-0.11	
-0.00	-0.15	0.23	-0.02	-0.12	-0.14	0.2
2	0.13	-0.02	-0.01	0.19	-0.20	
-0.19	-0.20	-0.07			1.00	
-0.17	-0.17	-0.46		-0.06		
Tax revenue (%)				-0.05	0.04	
-0.16	-0.20	-0.39	-0.15	-0.20	-0.24	-0.3
9	0.02	0.46	-0.12	-0.37	0.36	
-0.29		0.38	-0.19		-0.17	
1.00	-0.09	0.26		-0.18		
Total tax rate				-0.16	0.10	
0.17	0.17	0.17	0.12	0.10	0.31	0.18
0.09	0.06	0.10	0.18	-0.10		
0.14	-0.07	0.15			-0.17	
-0.09	1.00	0.03		0.18		
Unemployment rate				-0.12	0.13	
0.06	0.01	-0.03	0.02	0.18	0.16	-0.06
-0.11	-0.02	0.05	-0.02	-0.04		
0.01		0.08	-0.03		-0.46	
0.26	0.03	1.00		0.00		
Urban_population				-0.01	0.10	
0.54	0.89	-0.11	0.93	0.01	0.04	-0.10
-0.01	-0.08	0.78	-0.05	0.06		
0.07	-0.02	0.95			-0.06	
-0.18	0.18	0.00		1.00		

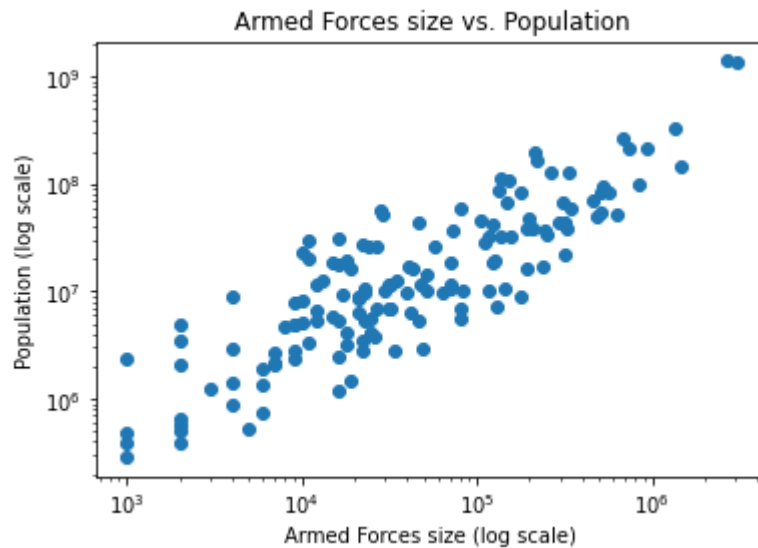


```
In [ ]: #scatter plot with a logarithmic scale
plt.scatter(x=df['Armed Forces size'], y=df['Population'])
plt.xscale('log')
plt.yscale('log')

# Labels
plt.xlabel('Armed Forces size (log scale)')
plt.ylabel('Population (log scale)')

plt.title('Armed Forces size vs. Population')
```

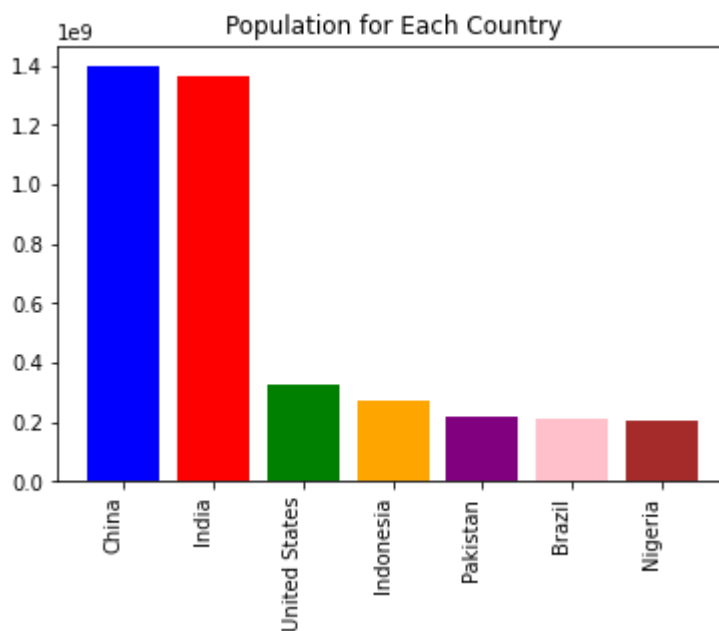
```
Out[ ]: Text(0.5, 1.0, 'Armed Forces size vs. Population')
```



Bar graph of top 7 most populated countries

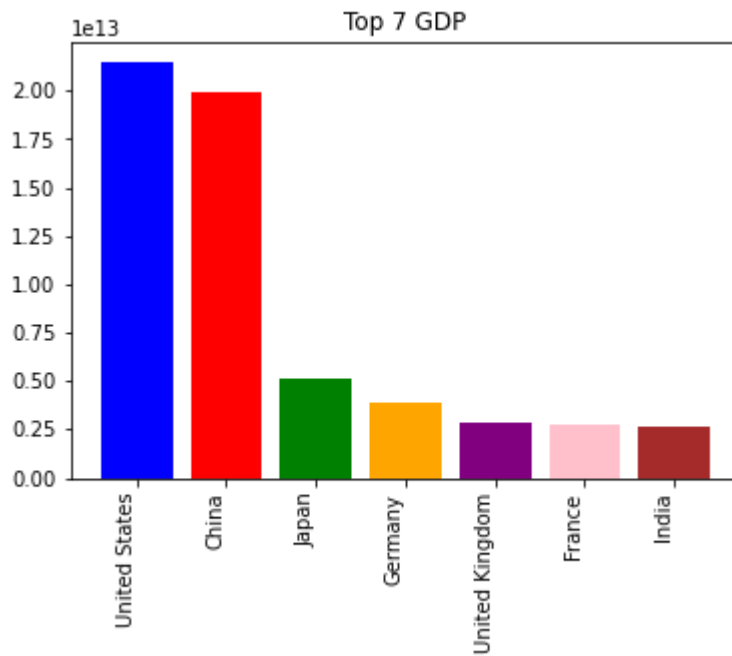
```
In [ ]: plt.bar(df.sort_values(by="Population", ascending=False).head(7)["Country"],
               df.sort_values(by="Population", ascending=False).head(7)["Population"],
               color=['blue', 'red', 'green', 'orange', 'purple', 'pink', 'brown'])

plt.xticks(rotation=90, ha='right')
plt.title('Population for Each Country')
plt.show()
```



```
In [ ]: plt.bar(df.sort_values(by="GDP", ascending=False).head(7)["Country"],
               df.sort_values(by="GDP", ascending=False).head(7)["GDP"],
               color=['blue', 'red', 'green', 'orange', 'purple', 'pink', 'brown'])

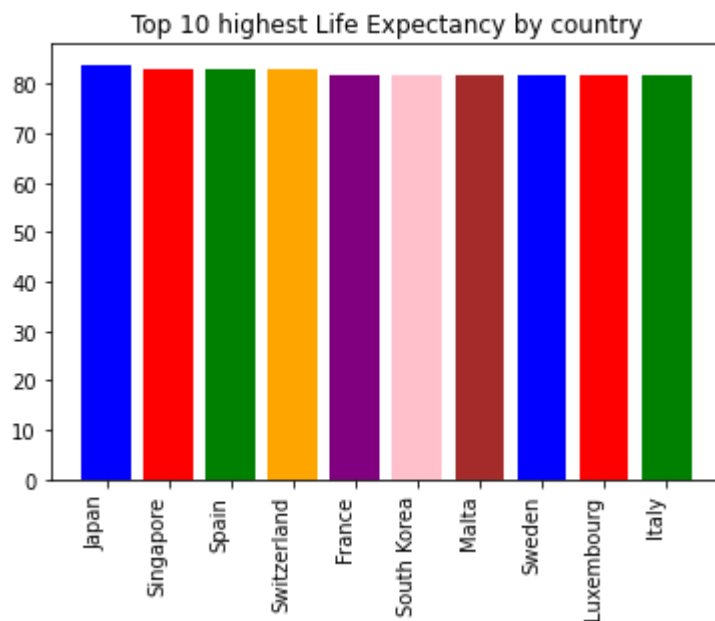
plt.xticks(rotation=90, ha='right')
plt.title('Top 7 GDP')
plt.show()
```



```
In [ ]: plt.bar(df.sort_values(by="Life expectancy",ascending=False).head(10)["Country"],
               df.sort_values(by="Life expectancy",ascending=False).head(10)["Life expectancy"]
               color=['blue','red','green','orange','purple','pink','brown'])

plt.xticks(rotation=90, ha='right')
plt.title('Top 10 highest Life Expectancy by country')

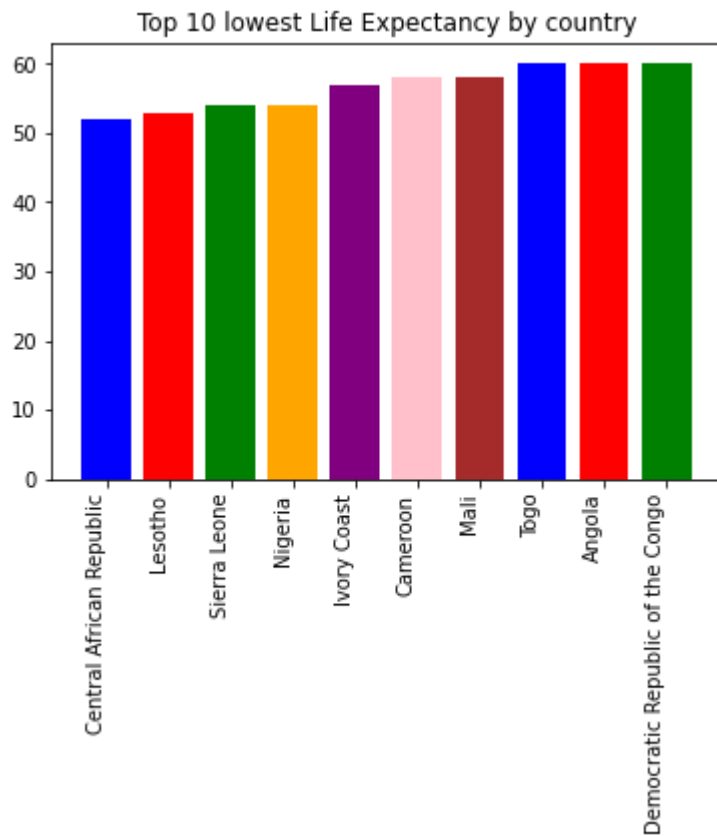
plt.show()
```



```
In [ ]: plt.bar(df.sort_values(by="Life expectancy",ascending=True).head(10)["Country"],
               df.sort_values(by="Life expectancy",ascending=True).head(10)["Life expectancy"]
               color=['blue','red','green','orange','purple','pink','brown'])

plt.xticks(rotation=90, ha='right')
plt.title('Top 10 lowest Life Expectancy by country')

plt.show()
```



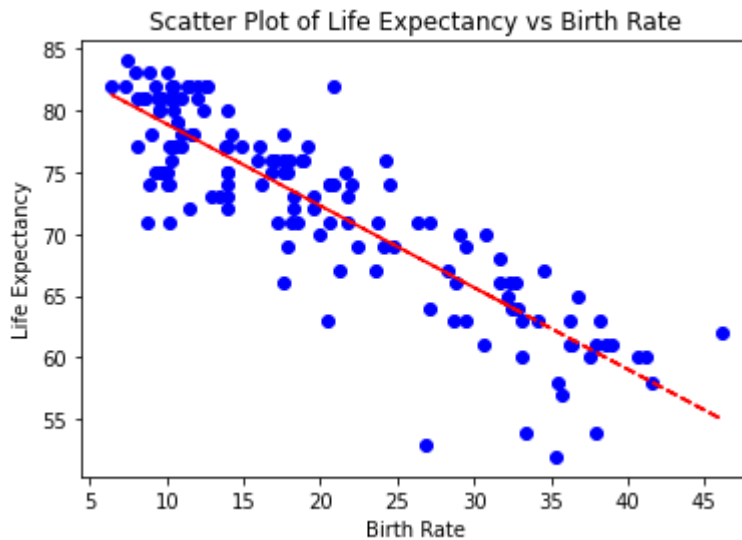
```
In [ ]: birth_rate = df['Birth Rate']
life_expectancy = df['Life expectancy']

plt.scatter(birth_rate, life_expectancy, color='blue')
plt.xlabel('Birth Rate')
plt.ylabel('Life Expectancy')
plt.title('Scatter Plot of Life Expectancy vs Birth Rate')

# Add best fit line

#get coefficients for best fit
coefficients = np.polyfit(birth_rate, life_expectancy, 1)
#equation of best fit line
polynomial = np.poly1d(coefficients)
plt.plot(birth_rate, polynomial(birth_rate), color='red', linestyle='--', label='Line o

plt.show()
```



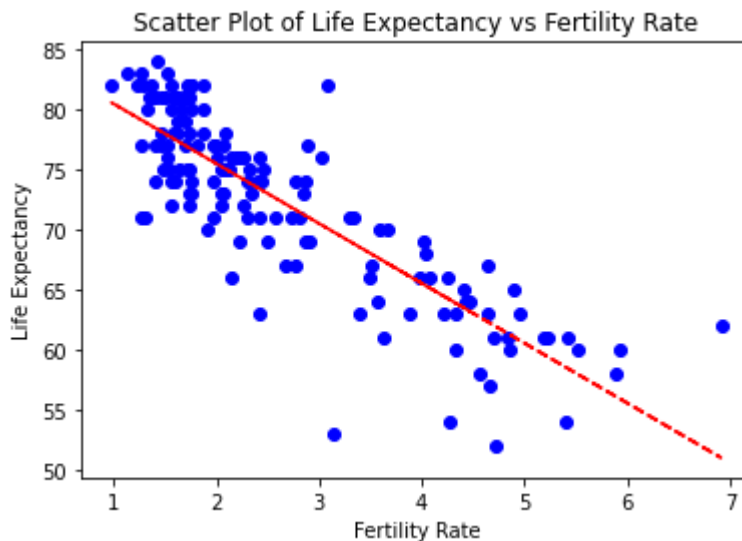
```
In [ ]: fertility_rate = df['Fertility Rate']
life_expectancy = df['Life expectancy']

plt.scatter(fertility_rate, life_expectancy, color='blue')
plt.xlabel('Fertility Rate')
plt.ylabel('Life Expectancy')
plt.title('Scatter Plot of Life Expectancy vs Fertility Rate')

# Add a line of best fit

coefficients = np.polyfit(fertility_rate, life_expectancy, 1)
polynomial = np.poly1d(coefficients)
plt.plot(fertility_rate, polynomial(fertility_rate), color='red', linestyle='--', label

plt.show()
```



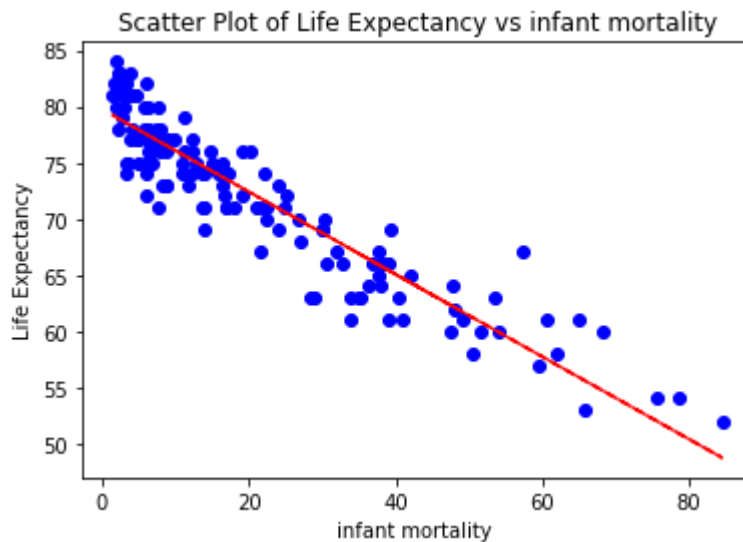
```
In [ ]: infant_mortality = df['Infant mortality']
life_expectancy = df['Life expectancy']

plt.scatter(infant_mortality, life_expectancy, color='blue')
plt.xlabel('infant mortality ')
```

```
plt.ylabel('Life Expectancy')
plt.title('Scatter Plot of Life Expectancy vs infant mortality')

# Add best fit line
coefficients = np.polyfit(infant_mortality, life_expectancy, 1)
polynomial = np.poly1d(coefficients)
plt.plot(infant_mortality, polynomial(infant_mortality), color='red', linestyle='--')

plt.show()
```

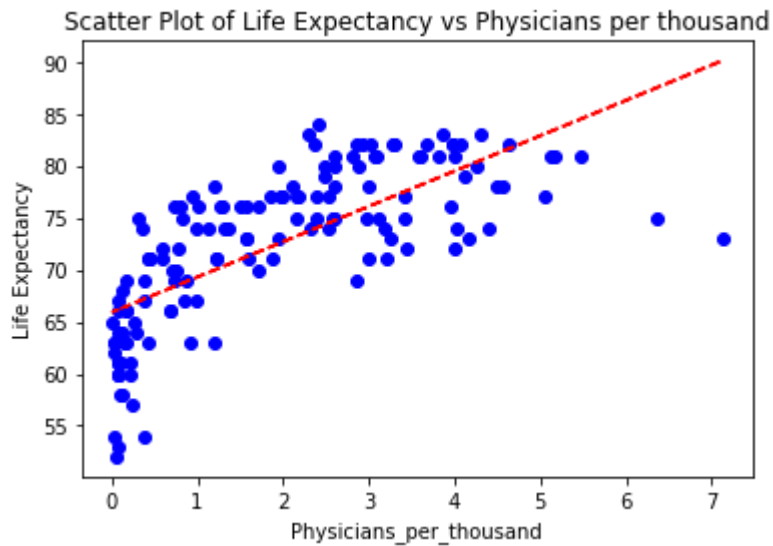


```
In [ ]: Physicians_per_thousand = df['Physicians per thousand']
life_expectancy = df['Life expectancy']

# Create a scatter plot
plt.scatter(Physicians_per_thousand, life_expectancy, color='blue')
plt.xlabel('Physicians_per_thousand')
plt.ylabel('Life Expectancy')
plt.title('Scatter Plot of Life Expectancy vs Physicians per thousand')

# Add a Line of best fit
coefficients = np.polyfit(Physicians_per_thousand, life_expectancy, 1)
polynomial = np.poly1d(coefficients)
plt.plot(Physicians_per_thousand, polynomial(Physicians_per_thousand), color='red', lin
```

```
Out[ ]: [<matplotlib.lines.Line2D at 0x2120e4a3970>]
```



EDA, Clustering

Hierarchical Clustering

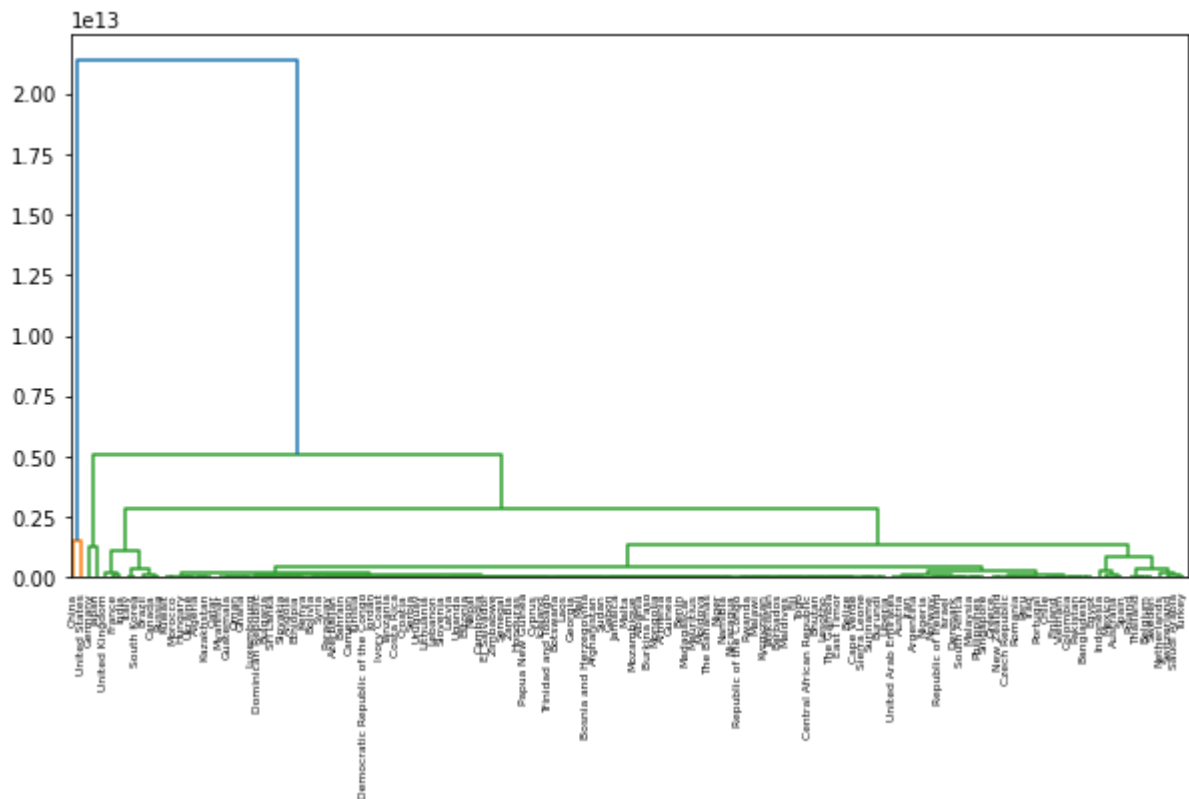
```
In [ ]: from scipy.cluster.hierarchy import linkage, dendrogram
import matplotlib.pyplot as plt
from sklearn.preprocessing import normalize
```

We need to separate Country col from the rest of the df

```
In [ ]: # remove the 'Country' column from the DataFrame before clustering
numerical_features_df = df.drop(columns=['Country'])
country=df.Country.values
```

```
In [ ]: mergings = linkage(numerical_features_df, method='complete')
plt.figure(figsize=(10, 5))

dendrogram(mergings,
            labels=country,
            leaf_rotation=90,
            leaf_font_size=6,
)
plt.show()
```



Feature Elimination

Recursive Feature Elimination Cross Validation to reduce features to find out the most important ones

```
In [ ]: import matplotlib.pyplot as plt
from sklearn.svm import SVC
from sklearn.model_selection import StratifiedKFold
from sklearn.feature_selection import RFECV
from sklearn.datasets import make_classification
%matplotlib inline
```

```
In [ ]: from sklearn.linear_model import LogisticRegression
#max_iter to make sure algorithm reaches convergence
logit=LogisticRegression(multi_class='ovr', solver='lbfgs', max_iter=200000)
```

Set X and y to features and target

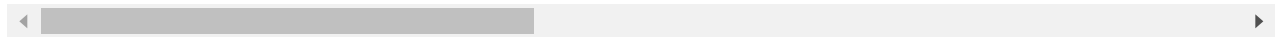
```
In [ ]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X=df.drop(columns=['Country', 'Life expectancy'])
y=df['Life expectancy']

X
```


Out []:

	Density\n(P/Km2)	Agricultural Land(%)	Land Area(Km2)	Armed Forces size	Birth Rate	Co2- Emissions	CPI	CPI Change (%)	Fertility Rate
0	60.00	58.10	652230.00	323000.00	32.49	8672.00	149.90	2.30	4.47
1	105.00	43.10	28748.00	9000.00	11.78	4536.00	119.05	1.40	1.62
2	18.00	17.40	2381741.00	317000.00	24.28	150006.00	151.36	2.00	3.02
4	26.00	47.50	1246700.00	117000.00	40.73	34693.00	261.73	17.10	5.52
6	17.00	54.30	2780400.00	105000.00	17.02	201348.00	232.75	53.50	2.26
...
186	36.00	44.40	9833517.00	1359000.00	11.60	5006302.00	117.24	7.50	1.73
187	20.00	82.60	176215.00	22000.00	13.86	6766.00	202.92	7.90	1.97
191	314.00	39.30	331210.00	522000.00	16.75	192668.00	163.52	2.80	2.05
193	25.00	32.10	752618.00	16000.00	36.19	5141.00	212.31	9.20	4.63
194	38.00	41.90	390757.00	51000.00	30.68	10983.00	105.51	0.90	3.62

146 rows × 21 columns



RFECV without scaling X (we get a warning if we dont scale the data, beacuse the algorithm is having difficulty reaching **convergence**)

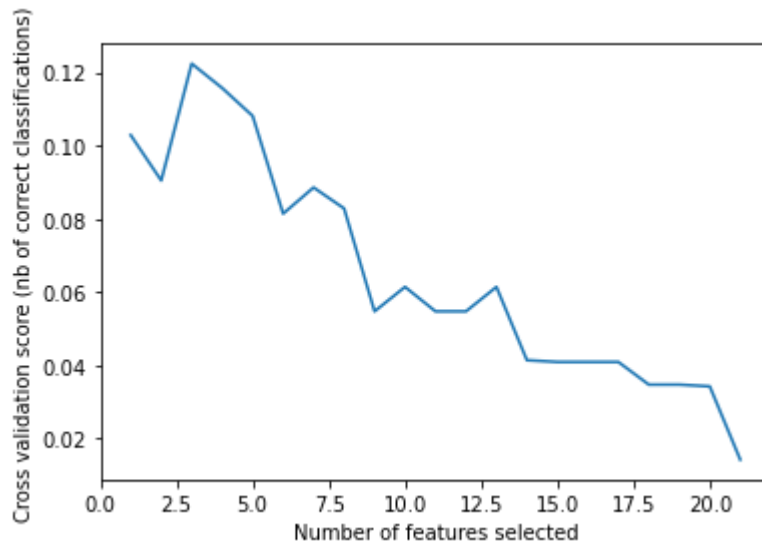
```
In [ ]: # Reduced number of splits and 'roc_auc' scoring
rfecv = RFECV(estimator=logit, step=1, cv=StratifiedKFold(10), scoring='accuracy')
rfecv.fit(X, y)
```

```
In [ ]: print(f"Optimal number of features : {rfecv.n_features_}")
```

Optimal number of features : 3

3 features yields the highest cross validation score

```
In [ ]: plt.figure()
plt.xlabel("Number of features selected")
plt.ylabel("Cross validation score (nb of correct classifications)")
plt.plot(range(1, len(rfecv.grid_scores_) + 1), rfecv.grid_scores_)
plt.show()
```



```
In [ ]: print(rfecv.support_)
        print("Selected Features: ",X.columns[rfecv.support_])

[False False False False  True False False False  True False False False
  False False  True False False False False False False]
Selected Features:  Index(['Birth Rate', 'Fertility Rate', 'Physicians per thousand'], d
type='object')
```

Ranking = 1 means that RFECV has chosen the feature

```
In [ ]: rfecv.ranking_
```

```
Out[ ]: array([12,  9, 15, 14,  1, 13, 11,  4,  1, 10, 18, 19,  2,  8,  1, 17,  5,
         6,  7,  3, 16])
```

RFECV with sacling X, to make sure the algorithm reaches convergence

```
In [ ]: from sklearn.preprocessing import StandardScaler
        scaler = StandardScaler()
        X_scaled = scaler.fit_transform(X)
        rfecv = RFECV(estimator=logit, step=1, cv=StratifiedKFold(10), scoring='accuracy')
        rfecv.fit(X_scaled, y)
```

c:\Users\almas\anaconda3\lib\site-packages\sklearn\model_selection_split.py:666: UserWarning: The least populated class in y has only 1 members, which is less than n_splits=10.

```
Out[ ]: warnings.warn(("The least populated class in y has only %d"
        RFECV(cv=StratifiedKFold(n_splits=10, random_state=None, shuffle=False),
        estimator=LogisticRegression(max_iter=200000, multi_class='ovr'),
        scoring='accuracy'))
```

```
In [ ]: print(f"Optimal number of features : {rfecv.n_features_}")
        selected_features= X.columns[rfecv.support_]
        print("Selected Features: ",selected_features)

        #create a df using the selected features
        selected_features_df = X[selected_features]
        selected_features_df
```

Optimal number of features : 16

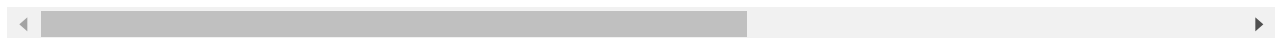
Selected Features: Index(['Agricultural Land(%)', 'Land Area(Km2)', 'Armed Forces size',

'Birth Rate', 'CPI', 'CPI Change (%)', 'Fertility Rate',
'Forested Area (%)', 'Gasoline Price', 'Infant mortality',
'Out of pocket health expenditure', 'Physicians per thousand',
'Population: Labor force participation (%)', 'Tax revenue (%)',
'Total tax rate', 'Unemployment rate'],
dtype='object')

Out[]:

	Agricultural Land(%)	Land Area(Km2)	Armed Forces size	Birth Rate	CPI	CPI Change (%)	Fertility Rate	Forested Area (%)	Gasoline Price	Infant mortality
0	58.10	652230.00	323000.00	32.49	149.90	2.30	4.47	2.10	0.70	47.90
1	43.10	28748.00	9000.00	11.78	119.05	1.40	1.62	28.10	1.36	7.80
2	17.40	2381741.00	317000.00	24.28	151.36	2.00	3.02	0.80	0.28	20.10
4	47.50	1246700.00	117000.00	40.73	261.73	17.10	5.52	46.30	0.97	51.60
6	54.30	2780400.00	105000.00	17.02	232.75	53.50	2.26	9.80	1.10	8.80
...
186	44.40	9833517.00	1359000.00	11.60	117.24	7.50	1.73	33.90	0.71	5.60
187	82.60	176215.00	22000.00	13.86	202.92	7.90	1.97	10.70	1.50	6.40
191	39.30	331210.00	522000.00	16.75	163.52	2.80	2.05	48.10	0.80	16.50
193	32.10	752618.00	16000.00	36.19	212.31	9.20	4.63	65.20	1.40	40.40
194	41.90	390757.00	51000.00	30.68	105.51	0.90	3.62	35.50	1.34	33.90

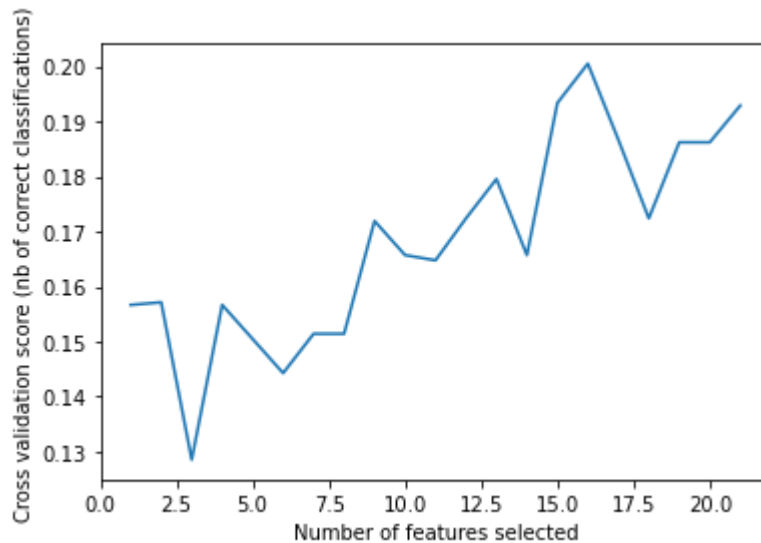
146 rows × 16 columns



16 features yields the highest cross validation score

In []:

```
#plot CV score with number of features selected
plt.figure()
plt.xlabel("Number of features selected")
plt.ylabel("Cross validation score (nb of correct classifications)")
plt.plot(range(1, len(rfecv.grid_scores_) + 1), rfecv.grid_scores_)
plt.show()
```



PCA

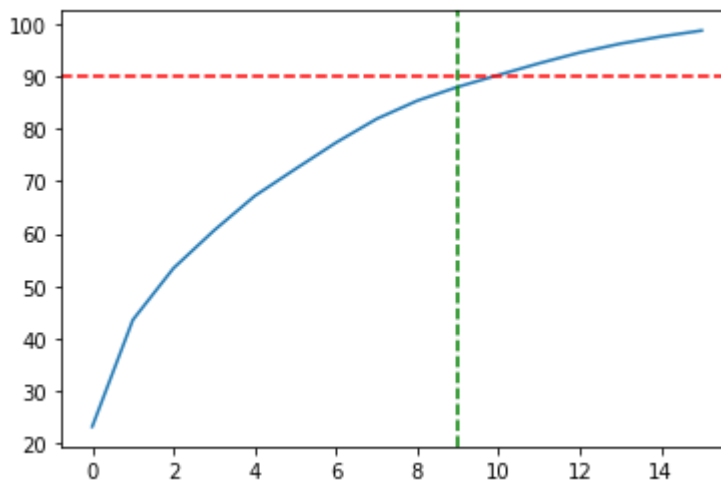
Dimensionality reduction algorithm on ***selected_features_df***, to reduce the dimensions of the dataframe

```
In [ ]: from sklearn.decomposition import PCA
pca_Object = PCA(n_components=16)
pca_Object.fit(X_scaled)
pca_X = pca_Object.transform(X)
```

```
In [ ]: #The amount of variance that each PC explains
var_Data= pca_Object.explained_variance_ratio_
#Cumulative Variance explains
var1_Data=np.cumsum(np.round(pca_Object.explained_variance_ratio_, decimals=4)*100)
var1_Data
```

```
Out[ ]: array([23.18, 43.54, 53.4 , 60.61, 67.18, 72.32, 77.35, 81.87, 85.34,
      87.98, 90.24, 92.47, 94.52, 96.21, 97.57, 98.71])
```

```
In [ ]: plt.plot(var1_Data)
plt.axhline(y=90, color='red', linestyle='--')
plt.axvline(x=9, color='green', linestyle='--')
plt.show()
```



Selecting features that have > 90% of the variance (9 features in this case)

```
In [ ]: from sklearn.decomposition import PCA
from sklearn import decomposition

pcaOBJ = decomposition.PCA(n_components=9)
pcaOBJ.fit(pca_X)
obj = pcaOBJ.transform(pca_X)
print("obj shape: ",obj.shape)
obj
```

```
Out[ ]: obj shape: (146, 9)
array([[ -4.87493786e+11,  7.29694423e+06, -9.48837518e+04, ...,
        -1.97292447e+02, -1.57562916e+01, -4.40227242e+01],
       [ -4.90576478e+11, -1.92718691e+07, -5.03016090e+05, ...,
        -1.34243206e+02, -2.79482985e+01, -9.88525772e+00],
       [ -3.65685046e+11,  1.47602551e+07,  1.98171091e+06, ...,
        -1.47821908e+02, -2.85436096e+01, -1.51489105e+01],
       ...,
       [ -2.91477272e+11,  4.78513204e+07, -6.13503980e+05, ...,
        1.94506550e+01, -8.08257122e+00,  1.49074173e+01],
       [ -4.84291735e+11, -7.01658472e+06,  1.50381783e+05, ...,
        -1.82729399e+02,  5.22851105e+01,  3.78757973e+01],
       [ -4.85602729e+11, -1.04848006e+07, -2.28985431e+05, ...,
        -1.91586820e+02, -5.09214783e+01,  6.35052490e+00]])
```

```
In [ ]: import pandas as pd

# Convert the 'obj' array into a DataFrame
pca_df = pd.DataFrame(obj)

# Print the DataFrame containing the chosen principal components with updated column names
print(pca_df)
```

	0	1	2	3	4 \
0	-487493786286.97	7296944.23	-94883.75	293155.97	25075.00
1	-490576478144.36	-19271869.06	-503016.09	-38965.69	-7349.40
2	-365685046327.28	14760255.08	1981710.91	-134068.24	19735.61
3	-426515000792.43	5499636.44	765716.05	-118533.82	-446.08
4	-139910852947.05	9360349.40	2375281.85	-338646.30	-10305.89
..
141	16794992860626.67	-536659176.46	-1665092.48	3011280.81	40423.10

```

142 -457665749541.99 -19774800.46 -351530.98 -51741.36 -5544.30
143 -291477272304.45 47851320.39 -613503.98 224789.70 40564.44
144 -484291734611.70 -7016584.72 150381.78 67663.12 -13902.30
145 -485602729205.39 -10484800.58 -228985.43 83918.12 -7096.85

```

```

      5      6      7      8
0 -10631.59 -197.29 -15.76 -44.02
1 -2708.75 -134.24 -27.95 -9.89
2 -391.50 -147.82 -28.54 -15.15
3 6531.91 -170.07 89.28 11.03
4 16199.47 -105.59 57.68 -30.45
..      ...      ...      ...      ...
141 20006.63 -177.45 58.16 -8.67
142 -1615.29 -215.07 44.31 -36.88
143 -15125.34 19.45 -8.08 14.91
144 2864.06 -182.73 52.29 37.88
145 -2157.83 -191.59 -50.92 6.35

```

[146 rows x 9 columns]

Split data into train and test (for both original dataframe and PCA dataframe)

Split data

```

In [ ]: #set X and Y
df.columns
df_X = df.drop(columns=['Life expectancy', 'Country'])
df_y=df['Life expectancy']
df_X

```

```

Out[ ]:

```

	Density\n(P/Km2)	Agricultural Land(%)	Land Area(Km2)	Armed Forces size	Birth Rate	Co2- Emissions	CPI	CPI Change (%)	Fertility Rate
0	60.00	58.10	652230.00	323000.00	32.49	8672.00	149.90	2.30	4.47
1	105.00	43.10	28748.00	9000.00	11.78	4536.00	119.05	1.40	1.62
2	18.00	17.40	2381741.00	317000.00	24.28	150006.00	151.36	2.00	3.02
4	26.00	47.50	1246700.00	117000.00	40.73	34693.00	261.73	17.10	5.52
6	17.00	54.30	2780400.00	105000.00	17.02	201348.00	232.75	53.50	2.26
...
186	36.00	44.40	9833517.00	1359000.00	11.60	5006302.00	117.24	7.50	1.73
187	20.00	82.60	176215.00	22000.00	13.86	6766.00	202.92	7.90	1.97
191	314.00	39.30	331210.00	522000.00	16.75	192668.00	163.52	2.80	2.05
193	25.00	32.10	752618.00	16000.00	36.19	5141.00	212.31	9.20	4.63
194	38.00	41.90	390757.00	51000.00	30.68	10983.00	105.51	0.90	3.62

146 rows × 21 columns

```
In [ ]: from sklearn.model_selection import train_test_split

#original df
df_X_train, df_X_test, df_y_train, df_y_test = train_test_split(df_X, df_y, test_size=0

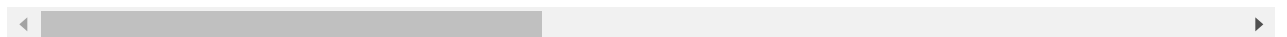
#PCA df
pca_X_train, pca_X_test, pca_y_train, pca_y_test = train_test_split(pca_df, df_y, test_
```

```
In [ ]: df_X_train
```

```
Out[ ]:
```

	Density\n(P/Km2)	Agricultural Land(%)	Land Area(Km2)	Armed Forces size	Birth Rate	Co2- Emissions	CPI	CPI Change (%)	Fertility Rate
119	3.00	47.10	824292.00	16000.00	28.64	4228.00	157.97	3.70	3.40
82	400.00	24.60	20770.00	178000.00	20.80	65166.00	108.15	0.80	3.09
27	463.00	79.20	27830.00	31000.00	39.01	495.00	182.11	-0.70	5.41
6	17.00	54.30	2780400.00	105000.00	17.02	201348.00	232.75	53.50	2.26
69	53.00	59.00	245857.00	13000.00	36.36	2996.00	262.95	9.50	4.70
...
81	72.00	64.50	70273.00	9000.00	12.50	37711.00	106.58	0.90	1.75
88	94.00	48.50	580367.00	29000.00	28.75	17910.00	180.51	4.70	3.49
102	203.00	61.40	118484.00	15000.00	34.12	1298.00	418.34	9.40	4.21
13	1265.00	70.60	148460.00	221000.00	18.18	84246.00	179.68	5.60	2.04
122	508.00	53.30	41543.00	41000.00	9.70	170780.00	115.91	2.60	1.59

102 rows × 21 columns



```
In [ ]: pca_X_train
```

```
Out[ ]:
```

	0	1	2	3	4	5	6	7	8
93	-492926879761.13	-19717685.69	256225.54	98781.60	-11653.80	-1115.61	-201.53	-0.06	-14.45
64	-183959199408.05	-27185565.61	-623347.54	-68353.35	18859.10	-6318.34	134.32	-43.21	-5.67
23	-500479450191.16	-13257115.23	-613016.44	93768.01	-8092.75	-3646.17	215.76	25.75	-31.44
4	-139910852947.05	9360349.40	2375281.85	-338646.30	-10305.89	16199.47	-105.59	57.68	-30.45
55	-491939923727.14	-11411869.72	-347960.33	30812.81	-10788.27	-372.50	-177.48	97.91	-16.82
...

	0	1	2	3	4	5	6	7	8
63	-189125853565.30	-32173212.42	-695049.22	66042.76	-7560.54	225.08	-172.24	-39.09	-25.86
70	-425818347608.60	15742854.47	-348130.07	326042.23	-22711.15	5840.07	-126.89	23.52	-17.46
81	-496722856174.43	-8073093.00	-583445.79	150327.81	-13331.71	-1680.22	-34.93	243.66	4.09
11	-258667874182.18	104258888.57	-1150506.61	265280.72	-14807.03	22589.15	995.19	37.40	-34.07
95	230953991129.07	-38835968.31	-872017.13	-73547.49	-3141.47	3172.27	254.38	-28.03	-19.89

102 rows × 9 columns

Regression learning algorithms

Linear Regression (Original DF)

```
In [ ]: # import model
        from sklearn.linear_model import LinearRegression

        # instantiate
        linreg = LinearRegression()

        # fit the model to the training data (learn the coefficients)
        linreg.fit(df_X_train, df_y_train)
```

Out[]: LinearRegression()

```
In [ ]: # print the intercept and coefficients
        print(linreg.intercept_)
        print(linreg.coef_)

85.4435090872766
[ 1.84055701e-04 -4.35098883e-02  3.47239592e-07  4.20157735e-06
  3.33298458e-02 -1.22994058e-05 -3.98249478e-03 -2.72815435e-02
 -1.26333909e+00 -4.05463409e-02  1.37548973e+00  2.60968148e-12
 -2.71481873e-01 -1.62726276e-02  3.41912389e-01  1.41881517e-08
 -5.08852319e-02 -2.35309539e-02  2.38153255e-02 -9.73249186e-02
 -2.19851545e-08]
```

```
In [ ]: # make predictions on the testing set
        df_y_pred = linreg.predict(df_X_test)
```

```
In [ ]: from sklearn import metrics
        df_lr_mae=metrics.mean_absolute_error(df_y_test, df_y_pred)
        df_lr_mse=metrics.mean_squared_error(df_y_test, df_y_pred)
        df_lr_rmse=np.sqrt(metrics.mean_squared_error(df_y_test, df_y_pred))

        print('Mean Absoulute Error',df_lr_mae)
        print('Mean Squared Error:',df_lr_mse)
        print('Root Mean Squared Error: ',df_lr_rmse)
```


Mean Absolute Error 3.2640263593722225
 Mean Squared Error: 73.34989071379162
 Root Mean Squared Error: 8.564455073954887

Linear Regression (PCA DF)

```
In [ ]: # instantiate
linreg = LinearRegression()

# fit the model to the training data (learn the coefficients)
linreg.fit(pca_X_train, pca_y_train)
```

Out[]: LinearRegression()

```
In [ ]: # print the intercept and coefficients
print(linreg.intercept_)
print(linreg.coef_)

71.60813416620944
[ 3.70777203e-13 -1.36824062e-08 -2.05564542e-07 -1.79670327e-09
 5.01305770e-05 2.22923475e-05 1.81176697e-03 -6.89317500e-02
 1.81977457e-02]
```

```
In [ ]: # make predictions on the testing set
pca_y_pred = linreg.predict(pca_X_test)
```

```
In [ ]: from sklearn import metrics
pca_lr_mae=metrics.mean_absolute_error(pca_y_test, pca_y_pred)
pca_lr_mse=metrics.mean_squared_error(pca_y_test, pca_y_pred)
pca_lr_rmse=np.sqrt(metrics.mean_squared_error(pca_y_test, pca_y_pred))

print('Mean Absolute Error',pca_lr_mae)
print('Mean Squared Error:',pca_lr_mse)
print('Root Mean Squared Error: ',pca_lr_rmse)
```

Mean Absolute Error 7.031591494863505
 Mean Squared Error: 169.2208982835762
 Root Mean Squared Error: 13.008493313354018

SVM (Original DF)

Without scaling

```
In [ ]: # Fitting SVR to the dataset
from sklearn.svm import SVR
regressor = SVR(kernel = 'rbf', gamma='auto')
regressor.fit(df_X_train, df_y_train)
```

Out[]: SVR(gamma='auto')

```
In [ ]: df_y_pred_no_scaling = regressor.predict(df_X_test)
```

```
In [ ]: predictions = [[206,43.2,301340,347000,7.3,320411,110.62,0.6,1.29,31.8,1.61,20012443920

# Scale the prediction data
scaler_X = StandardScaler()
scaler_X.fit(predictions)
scaled_prediction = scaler_X.transform(predictions)

# Predict using the scaled input
predicted_value = regressor.predict(scaled_prediction)

# Print the predicted value
print("Predicted Life expectancy:", predicted_value)
print('Actual Life expectancy: ',)
```

Predicted Life expectancy: [74.04545455]
 Actual Life expectancy: 82
 [74.04545455]
 Actual Life expectancy: 82

```
In [ ]: print('Mean Absolute Error',metrics.mean_absolute_error(df_y_test, df_y_pred_no_scaling)
print('Mean Squared Error:',metrics.mean_squared_error(df_y_test, df_y_pred_no_scaling)
print('Root Mean Squared Error: ',np.sqrt(metrics.mean_squared_error(df_y_test, df_y_pr
```

Mean Absolute Error 5.890495867768595
 Mean Squared Error: 55.81611570247933
 Root Mean Squared Error: 7.471018384563067

SVM with Scaling

```
In [ ]: #scaling
from sklearn.preprocessing import StandardScaler
scaler_X = StandardScaler()
df_X_scaled = scaler_X.fit_transform(df_X_train)
```

```
In [ ]: #create and fit SVR using scaled values
regressor = SVR(kernel='rbf', gamma='auto')
regressor.fit(X_scaled, y)
```

Out[]: SVR(gamma='auto')

```
In [ ]: predictions = [[206,43.2,301340,347000,7.3,320411,110.62,0.6,1.29,31.8,1.61,20012443920

# Scale the prediction data
scaled_prediction = scaler_X.transform(predictions)

# Predict using scaled input
predicted_value = regressor.predict(scaled_prediction)

# Print the predicted value
print("Predicted Life expectancy:", predicted_value)
print('Actual Life expectancy: 82')
```

Predicted Life expectancy: [80.88365092]
 Actual Life expectancy: 82

```
In [ ]: print('Mean Absolute Error',metrics.mean_absolute_error(df_y_test, df_y_pred))
        print('Mean Squared Error:',metrics.mean_squared_error(df_y_test, df_y_pred))
        print('Root Mean Squared Error: ',np.sqrt(metrics.mean_squared_error(df_y_test, df_y_pr
```

Mean Absolute Error 3.2640263593722225
Mean Squared Error: 73.34989071379162
Root Mean Squared Error: 8.564455073954887

SVM (PCA df)

Without scaling

```
In [ ]: regressor = SVR(kernel = 'rbf',gamma='auto')
        regressor.fit(pca_X_train, pca_y_train)
```

```
Out[ ]: SVR(gamma='auto')
```

```
In [ ]: pca_y_pred_no_scaling = regressor.predict(pca_X_test)
```

```
In [ ]: pca_predictions = [[-258667874182.18,104258888.57,-1150506.61,265280.72,-14807.03,22589

# Scale the prediction data
scaler_X = StandardScaler()
scaler_X.fit(pca_X_train)
pca_scaled_prediction = scaler_X.transform(pca_predictions)

# Predict using the scaled input
pca_predicted_value = regressor.predict(pca_scaled_prediction)

# Print the predicted value
print("Predicted Life expectancy:", pca_predicted_value)
print('Actual Life expectancy: 82')
```

Predicted Life expectancy: [74.04545455]
Actual Life expectancy: 82

```
In [ ]: pca_svm_mae=metrics.mean_absolute_error(pca_y_test, pca_y_pred)
        pca_svm_mse=metrics.mean_squared_error(pca_y_test, pca_y_pred)
        pca_svm_rmse=np.sqrt(metrics.mean_squared_error(pca_y_test, pca_y_pred))
        print('Mean Absolute Error',pca_svm_mae)
        print('Mean Squared Error:',pca_svm_mse)
        print('Root Mean Squared Error: ',pca_svm_rmse)
```

Mean Absolute Error 7.031591494863505
Mean Squared Error: 169.2208982835762
Root Mean Squared Error: 13.008493313354018

With scaling

```
In [ ]: #scaling
        from sklearn.preprocessing import StandardScaler
        scaler_X = StandardScaler()
        pca_X_scaled = scaler_X.fit_transform(pca_X_train)
```

```
In [ ]: #create and fit SVR using scaled values
regressor = SVR(kernel='rbf', gamma='auto')
regressor.fit(pca_X_scaled, pca_y_train)
```

```
Out[ ]: SVR(gamma='auto')
```

```
In [ ]: pca_predictions = [[-258667874182.18,104258888.57,-1150506.61,265280.72,-14807.03,22589

# Scale the prediction data
scaler_X = StandardScaler()
scaler_X.fit(pca_X_train)
pca_scaled_prediction = scaler_X.transform(pca_predictions)

# Predict using the scaled input
pca_predicted_value = regressor.predict(pca_scaled_prediction)

# Print the predicted value
print("Predicted Life expectancy:", pca_predicted_value)
print('Actual Life expectancy: 80')
```

Predicted Life expectancy: [71.62540858]
Actual Life expectancy: 80

```
In [ ]: pca_svmS_mae=metrics.mean_absolute_error(pca_y_test, pca_y_pred)
pca_svmS_mse=metrics.mean_squared_error(pca_y_test, pca_y_pred)
pca_svmS_rmse=np.sqrt(metrics.mean_squared_error(pca_y_test, pca_y_pred))
print('Mean Absolute Error',pca_svmS_mae)
print('Mean Squared Error:',pca_svmS_mse)
print('Root Mean Squared Error: ',pca_svmS_rmse)
```

Mean Absolute Error 7.031591494863505
Mean Squared Error: 169.2208982835762
Root Mean Squared Error: 13.008493313354018

KNN Regression (Original DF)

```
In [ ]: from sklearn.neighbors import KNeighborsRegressor
from sklearn.model_selection import cross_val_score

knn_regressor = KNeighborsRegressor(n_neighbors=5)
#Use scaled X
knn_regressor.fit(df_X_train, df_y_train)
```

```
Out[ ]: KNeighborsRegressor()
```

```
In [ ]: from sklearn.model_selection import cross_val_score
from sklearn.neighbors import KNeighborsRegressor

k_range = list(range(1, 31))
k_scores_MAE = []
k_scores_MSE = []
k_scores_RMSE = []
min_k_MAE = k_range[0]
min_k_MSE = k_range[0]
min_k_RMSE = k_range[0]
```

```

for k in k_range:
    knn = KNeighborsRegressor(n_neighbors=k)
    # Use 'neg_mean_absolute_error' as the scoring metric
    scores_MAE = cross_val_score(knn, df_X_train, df_y_train, cv=10, scoring='neg_mean_
    k_scores_MAE.append(-scores_MAE.mean())

for k in k_range:
    knn = KNeighborsRegressor(n_neighbors=k)
    # Use 'neg_mean_squared_error' as the scoring metric
    scores_MSE = cross_val_score(knn, df_X_train, df_y_train, cv=10, scoring='neg_mean_
    k_scores_MSE.append(-scores_MSE.mean())

for k in k_range:
    knn = KNeighborsRegressor(n_neighbors=k)
    # Use 'neg_root_mean_squared_error' as the scoring metric
    scores_RMSE = cross_val_score(knn, df_X_train, df_y_train, cv=10, scoring='neg_root
    k_scores_RMSE.append(-scores_RMSE.mean())

k_scores_MAE.sort()
k_scores_MSE.sort()
k_scores_RMSE.sort()

print('Minimum MAE:', k_scores_MAE[0], 'at k =', min_k_MAE)
print('Minimum MSE:', k_scores_MSE[0], 'at k =', min_k_MSE)
print('Minimum RMSE:', k_scores_RMSE[0], 'at k =', min_k_RMSE)

```

Minimum MAE: 5.559870129870131 at k = 1
 Minimum MSE: 46.58687899402185 at k = 1
 Minimum RMSE: 6.604713443110066 at k = 1

KNN Regression (PCA df)

```

In [ ]: from sklearn.neighbors import KNeighborsRegressor
        from sklearn.model_selection import cross_val_score

        knn_regressor = KNeighborsRegressor(n_neighbors=5)
        knn_regressor.fit(pca_X_train, pca_y_train)

```

Out[]: KNeighborsRegressor()

```

In [ ]: from sklearn.model_selection import cross_val_score
        from sklearn.neighbors import KNeighborsRegressor

        k_range = list(range(1, 31))
        pca_k_scores_MAE = []
        pca_k_scores_MSE = []
        pca_k_scores_RMSE = []
        min_k_MAE = k_range[0]
        min_k_MSE = k_range[0]
        min_k_RMSE = k_range[0]
        for k in k_range:
            knn = KNeighborsRegressor(n_neighbors=k)
            # Use 'neg_mean_absolute_error' as the scoring metric
            scores_MAE = cross_val_score(knn, pca_X_train, pca_y_train, cv=10, scoring='neg_me
            pca_k_scores_MAE.append(-scores_MAE.mean())

```

```

for k in k_range:
    knn = KNeighborsRegressor(n_neighbors=k)
    # Use 'neg_mean_squared_error' as the scoring metric
    scores_MSE = cross_val_score(knn, pca_X_train, pca_y_train, cv=10, scoring='neg_mea
    pca_k_scores_MSE.append(-scores_MSE.mean())

for k in k_range:
    knn = KNeighborsRegressor(n_neighbors=k)
    # Use 'neg_root_mean_squared_error' as the scoring metric
    scores_RMSE = cross_val_score(knn, pca_X_train, pca_y_train, cv=10, scoring='neg_ro
    pca_k_scores_RMSE.append(-scores_RMSE.mean())

pca_k_scores_MAE.sort()
pca_k_scores_MSE.sort()
pca_k_scores_RMSE.sort()

print("MAE scores:")
for k, score in zip(k_range, k_scores_MAE):
    print("k =", k, "MAE =", score)

print("\nMSE scores:")
for k, score in zip(k_range, k_scores_MSE):
    print("k =", k, "MSE =", score)

print("\nRMSE scores:")
for k, score in zip(k_range, k_scores_RMSE):
    print("k =", k, "RMSE =", score)

```

MAE scores:

```

k = 1 MAE = 5.559870129870131
k = 2 MAE = 5.577090909090909
k = 3 MAE = 5.58603305785124
k = 4 MAE = 5.5901010101010105
k = 5 MAE = 5.60267942583732
k = 6 MAE = 5.618636363636363
k = 7 MAE = 5.630121212121212
k = 8 MAE = 5.634743083003952
k = 9 MAE = 5.646420454545455
k = 10 MAE = 5.6551636363636355
k = 11 MAE = 5.65979020979021
k = 12 MAE = 5.661643356643357
k = 13 MAE = 5.661711229946524
k = 14 MAE = 5.671598746081505
k = 15 MAE = 5.685422077922078
k = 16 MAE = 5.687077922077922
k = 17 MAE = 5.688212121212121
k = 18 MAE = 5.690033670033669
k = 19 MAE = 5.7524999999999995
k = 20 MAE = 5.7841414141414145
k = 21 MAE = 5.797909090909091
k = 22 MAE = 5.828512396694215
k = 23 MAE = 5.880795454545455
k = 24 MAE = 5.945454545454545
k = 25 MAE = 5.991363636363635
k = 26 MAE = 6.083090909090909
k = 27 MAE = 6.287954545454546
k = 28 MAE = 6.824545454545455
k = 29 MAE = 7.129545454545455
k = 30 MAE = 7.422727272727272

```

MSE scores:

k = 1 MSE = 46.58687899402185
k = 2 MSE = 46.69948181818183
k = 3 MSE = 46.777253787878784
k = 4 MSE = 46.85133358377159
k = 5 MSE = 46.86868446654417
k = 6 MSE = 47.07009548144164
k = 7 MSE = 47.081327999999999
k = 8 MSE = 47.1026393939394
k = 9 MSE = 47.14599345253086
k = 10 MSE = 47.21397379406308
k = 11 MSE = 47.361157251527615
k = 12 MSE = 47.44026936026937
k = 13 MSE = 47.47800481182334
k = 14 MSE = 47.91210443535704
k = 15 MSE = 48.136615767045456
k = 16 MSE = 48.20710707070707
k = 17 MSE = 48.71043033889187
k = 18 MSE = 49.41036641929499
k = 19 MSE = 49.82027146464647
k = 20 MSE = 50.300681818181815
k = 21 MSE = 50.63913580246914
k = 22 MSE = 51.09359128474831
k = 23 MSE = 51.74509943181819
k = 24 MSE = 52.9969573283859
k = 25 MSE = 54.05542929292928
k = 26 MSE = 57.44425454545454
k = 27 MSE = 61.12335227272727
k = 28 MSE = 69.9249494949495
k = 29 MSE = 76.77386363636363
k = 30 MSE = 90.9009090909091

RMSE scores:

k = 1 RMSE = 6.604713443110066
k = 2 RMSE = 6.6177688685869995
k = 3 RMSE = 6.621117640262172
k = 4 RMSE = 6.626885443241572
k = 5 RMSE = 6.649646223739947
k = 6 RMSE = 6.649748701978079
k = 7 RMSE = 6.653014433461443
k = 8 RMSE = 6.6591523090502465
k = 9 RMSE = 6.668777816434617
k = 10 RMSE = 6.670497428290341
k = 11 RMSE = 6.676416612257256
k = 12 RMSE = 6.678124672699113
k = 13 RMSE = 6.685631333103286
k = 14 RMSE = 6.720406837360675
k = 15 RMSE = 6.748622374986992
k = 16 RMSE = 6.760062472606554
k = 17 RMSE = 6.787917913681817
k = 18 RMSE = 6.83219262360767
k = 19 RMSE = 6.869665695377165
k = 20 RMSE = 6.887678949578192
k = 21 RMSE = 6.888104366033588
k = 22 RMSE = 6.951764087690283
k = 23 RMSE = 6.952990742983587
k = 24 RMSE = 7.053922891738419
k = 25 RMSE = 7.111311505688965
k = 26 RMSE = 7.320246818563701

```
k = 27 RMSE = 7.5610146161599605
k = 28 RMSE = 8.098451858996993
k = 29 RMSE = 8.405167919718448
k = 30 RMSE = 9.098034490730683
```

Decision Tree (Original DF)

```
In [ ]: from sklearn.tree import DecisionTreeClassifier
```

```
In [ ]: dtree = DecisionTreeClassifier()
```

```
In [ ]: dtree.fit(df_X_train,df_y_train)
```

```
Out[ ]: DecisionTreeClassifier()
```

```
In [ ]: df_y_pred_dt = dtree.predict(df_X_test)
```

```
In [ ]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score, explained_variance_score

df_dt_mae = mean_absolute_error(df_y_test, df_y_pred_dt)
df_dt_mse = mean_squared_error(df_y_test, df_y_pred_dt)
df_dt_rmse = np.sqrt(df_dt_mse)
df_dt_r2 = r2_score(df_y_test, df_y_pred_dt)
df_dt_ev = explained_variance_score(df_y_test, df_y_pred_dt)
print("Mean Absolute Error:", df_dt_mae)
print("Mean Squared Error:", df_dt_mse)
print("Root Mean Squared Error:", df_dt_rmse)
print("R-squared:", df_dt_r2)
print("Explained Variance Score:", df_dt_ev)
```

```
Mean Absolute Error: 3.0
Mean Squared Error: 16.045454545454547
Root Mean Squared Error: 4.005677788521506
R-squared: 0.7042509639643928
Explained Variance Score: 0.7088589517779789
```

Decision Tree (PCA DF)

```
In [ ]: from sklearn.tree import DecisionTreeClassifier
```

```
In [ ]: dtree = DecisionTreeClassifier()
```

```
In [ ]: dtree.fit(pca_X_train,pca_y_train)
```

```
Out[ ]: DecisionTreeClassifier()
```

```
In [ ]: pca_y_pred_dt = dtree.predict(pca_X_test)
```



```
In [ ]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score, explained_variance_score

pca_dt_mae = mean_absolute_error(pca_y_test, pca_y_pred_dt)
pca_dt_mse = mean_squared_error(pca_y_test, pca_y_pred_dt)
pca_dt_rmse = np.sqrt(pca_dt_mse)
pca_dt_r2 = r2_score(pca_y_test, pca_y_pred_dt)
pca_dt_ev = explained_variance_score(pca_y_test, pca_y_pred_dt)
print("Mean Absolute Error:", pca_dt_mae)
print("Mean Squared Error:", pca_dt_mse)
print("Root Mean Squared Error:", pca_dt_rmse)
print("R-squared:", pca_dt_r2)
print("Explained Variance Score:", pca_dt_ev)
```

Mean Absolute Error: 4.5
Mean Squared Error: 37.09090909090909
Root Mean Squared Error: 6.0902306270706275
R-squared: 0.3163421716570667
Explained Variance Score: 0.3578140619793402

Random Forest (orgainal DF)

```
In [ ]: from sklearn.ensemble import RandomForestClassifier
```

```
In [ ]: rfc = RandomForestClassifier(n_estimators=600)
```

```
In [ ]: rfc.fit(df_X_train, df_y_train)
```

```
Out[ ]: RandomForestClassifier(n_estimators=600)
```

```
In [ ]: predictions_rf = rfc.predict(df_X_test)
```

```
In [ ]: from sklearn.metrics import classification_report, confusion_matrix
```

```
In [ ]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score, explained_variance_score

mae = mean_absolute_error(df_y_test, predictions_rf)
mse = mean_squared_error(df_y_test, predictions_rf)
rmse = np.sqrt(mse)
r2 = r2_score(df_y_test, predictions_rf)
ev = explained_variance_score(df_y_test, predictions_rf)
print("Mean Absolute Error:", mae)
print("Mean Squared Error:", mse)
print("Root Mean Squared Error:", rmse)
print("R-squared:", r2)
print("Explained Variance Score:", ev)
```

Mean Absolute Error: 2.0454545454545454
Mean Squared Error: 7.863636363636363
Root Mean Squared Error: 2.8042176027613057
R-squared: 0.8550578378635693
Explained Variance Score: 0.8564288094444709

Random Forest (PCA DF)

```
In [ ]: from sklearn.ensemble import RandomForestClassifier
```

```
In [ ]: rfc = RandomForestClassifier(n_estimators=600)
```

```
In [ ]: rfc.fit(pca_X_train,pca_y_train)
```

```
Out[ ]: RandomForestClassifier(n_estimators=600)
```

```
In [ ]: predictions_rf = rfc.predict(pca_X_test)
```

```
In [ ]: from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score, explained_variance_score

mae = mean_absolute_error(pca_y_test, predictions_rf)
mse = mean_squared_error(pca_y_test, predictions_rf)
rmse = np.sqrt(mse)
r2 = r2_score(pca_y_test, predictions_rf)
ev = explained_variance_score(pca_y_test, predictions_rf)
print("Mean Absolute Error:", mae)
print("Mean Squared Error:", mse)
print("Root Mean Squared Error:", rmse)
print("R-squared:", r2)
print("Explained Variance Score:", ev)
```

```
Mean Absolute Error: 3.7045454545454546
Mean Squared Error: 25.113636363636363
Root Mean Squared Error: 5.011350752405619
R-squared: 0.5371066787261389
Explained Variance Score: 0.543057076212691
```

Regression learning algorithms Summary

```
In [ ]: import pandas as pd
summary = {
    'Model': ['Linear Regression (ODF)', 'Linear Regression (PCA DF)', 'SVM (ODF NO SCA',
              'SVM (PCA DF With SCALING)', 'KNN (ODF)', 'KNN (PCA DF)', 'DT (ODF)', 'DT',
    'MAE': [3.26, 7.03, 5.89, 3.26, 7.03, 7.03, 5.56, 5.56, 3.00, 4.50, 2.05, 3.70],
    'MSE': [73.35, 169.22, 55.82, 73.35, 169.22, 169.22, 46.59, 46.78, 16.05, 37.09, 7.
    'RMSE': [8.56, 13.01, 7.47, 8.56, 13.01, 13.01, 6.60, 6.62, 4.01, 6.09, 2.80, 5.01]
}

df = pd.DataFrame(summary)
df
```

```
Out[ ]:      Model  MAE  MSE  RMSE
```

```
0      Linear Regression (ODF)  3.26  73.35  8.56
```

1	Model	MAE	MSE	RMSE
2	Linear Regression (PCA DF)	7.03	169.22	13.01
3	SVM (ODF NO SCALING)	5.89	55.82	7.47
4	SVM (ODF With SCALING)	3.26	73.35	8.56
5	SVM (PCA DF NO SCALING)	7.03	169.22	13.01
6	SVM (PCA DF With SCALING)	7.03	169.22	13.01
7	KNN (ODF)	5.56	46.59	6.60
8	KNN (PCA DF)	5.56	46.78	6.62
9	DT (ODF)	3.00	16.05	4.01
10	DT (PCA DF)	4.50	37.09	6.09
11	Random Forest (ODF)	2.05	7.86	2.80
12	Random Forest (PCA DF)	3.70	25.11	5.01

Report

Top 2 algorithms

Random Forest (Original DF):

The Random Forest algorithm applied to the original DataFrame demonstrates outstanding performance in predicting life expectancy. With a Mean Absolute Error (MAE) of 2.05 and a Root Mean Squared Error (RMSE) of 2.80, the model provides highly accurate predictions. The low MAE and RMSE scores indicate that the model's predictions are, on average, very close to the actual life expectancy values. The strength of Random Forest lies in its ability to handle complex relationships within the data, making it robust against overfitting. By combining multiple decision trees and averaging their outputs, Random Forest reduces the risk of individual tree biases and increases overall prediction accuracy. This makes it a powerful and reliable choice for life expectancy prediction.

Decision Tree (Original DF):

The Decision Tree algorithm, when applied to the original DataFrame, also exhibits commendable performance in predicting life expectancy. With a Mean Absolute Error (MAE) of 3.0 and a Root Mean Squared Error (RMSE) of 4.01, the model delivers relatively accurate predictions. Decision Trees are easy to understand and interpret, making them valuable for gaining insights into feature importance and the decision-making process. However, compared to Random Forest, Decision Trees might be more prone to overfitting, especially on complex datasets. Nonetheless, the model's performance is still satisfactory, and its simplicity and interpretability make it an attractive option for scenarios where model interpretability is of extreme importance or when dealing with smaller datasets.

The rest of the algorithms

The remaining algorithms, including Linear Regression and SVM, show varying levels of performance in predicting life expectancy. While SVM demonstrates good performance with proper scaling, Linear Regression performs reasonably well. However, both PCA-based models (Linear Regression and SVM) show lower accuracy compared to the original DataFrame. KNN models provide moderate performance but are sensitive to the choice of the number of

neighbors (k). In summary, SVM and Linear Regression models can provide useful insights, but **Random Forest and Decision Tree** models offer superior accuracy for life expectancy prediction.

In conclusion, both the Random Forest and Decision Tree models applied to the original DataFrame show promise in predicting life expectancy. The Random Forest stands out as the top performer, providing superior accuracy and robustness to complex data relationships. On the other hand, the Decision Tree offers simplicity and interpretability, making it suitable for scenarios where model transparency and explainability are essential. The choice between these models would depend on the specific requirements of the project and the trade-offs between accuracy and interpretability.

Correlation analysis

Birth Rate (-0.88 correlation):The strong negative correlation between life expectancy and birth rate indicates that countries with higher birth rates tend to have lower life expectancies. High birth rates can put a strain on healthcare systems and resources, impacting overall public health and access to medical care.

Fertility Rate (-0.85 correlation):The negative correlation between life expectancy and fertility rate suggests that countries with higher fertility rates also tend to have lower life expectancies. High fertility rates can lead to challenges in providing adequate healthcare and social services, which can impact population health and life expectancy.

Infant Mortality (-0.93 correlation):The significant negative correlation between life expectancy and infant mortality underscores the critical link between early-life health and overall life expectancy. Lower infant mortality rates indicate better access to healthcare and improved maternal and child health.

Physicians per Thousand (0.72 correlation):The positive correlation between life expectancy and the number of physicians per thousand reflects the importance of healthcare access and medical resources in improving life expectancy. Countries with more physicians per thousand are better equipped to provide medical care and preventive services, positively impacting life expectancy.

Salman Almaskati - 500922635